



# Introduction & Neural Networks

Intelligent Architectures (5LIL0)

dr Alexios Balatsoukas-Stimming

Department of Electrical Engineering, Electronic Systems Group

**Efficient deep learning algorithms and hardware architectures.**

# “Efficient...”

Efficiency

## Efficiency

The ratio of useful output  $O$  to total input  $I$ :

$$E = \frac{O}{I}$$

## Efficiency

The ratio of useful output  $O$  to total input  $I$ :

$$E = \frac{O}{I}$$

## Efficiency examples:

## Efficiency

The ratio of useful output  $O$  to total input  $I$ :

$$E = \frac{O}{I}$$

### Efficiency examples:

1. Automotive: km/L
2. Lighting: lm/W
3. Communications: bps/Hz
4. Computing: pJ/operation

# Is Improved Efficiency Always Beneficial?

## Jevons Paradox (William S. Jevons, 1865)

Increases in efficiency with which a resource is used tend to increase the rate of consumption of that resource due to increasing demand.

# Is Improved Efficiency Always Beneficial?

## Jevons Paradox (William S. Jevons, 1865)

Increases in efficiency with which a resource is used tend to increase the rate of consumption of that resource due to increasing demand.

### Examples:

1. Coal use (Jevons' original observation)

# Is Improved Efficiency Always Beneficial?

## Jevons Paradox (William S. Jevons, 1865)

Increases in efficiency with which a resource is used tend to increase the rate of consumption of that resource due to increasing demand.

### Examples:

1. Coal use (Jevons' original observation)
2. LED lights

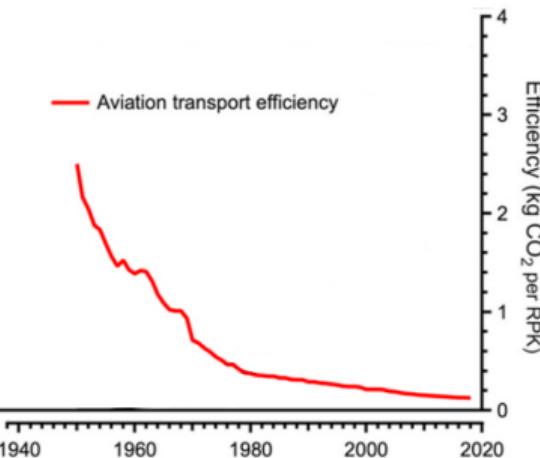
# Is Improved Efficiency Always Beneficial?

## Jevons Paradox (William S. Jevons, 1865)

Increases in efficiency with which a resource is used tend to increase the rate of consumption of that resource due to increasing demand.

### Examples:

1. Coal use (Jevons' original observation)
2. LED lights
3. Air travel [1]



[1] D. S. Lee *et al.*, "The contribution of global aviation to anthropogenic climate forcing for 2000 to 2018," 2021.

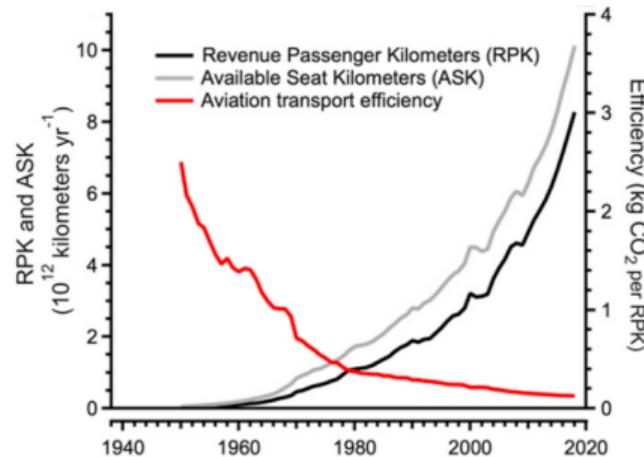
# Is Improved Efficiency Always Beneficial?

## Jevons Paradox (William S. Jevons, 1865)

Increases in efficiency with which a resource is used tend to increase the rate of consumption of that resource due to increasing demand.

### Examples:

1. Coal use (Jevons' original observation)
2. LED lights
3. Air travel [1]



[1] D. S. Lee *et al.*, "The contribution of global aviation to anthropogenic climate forcing for 2000 to 2018," 2021.

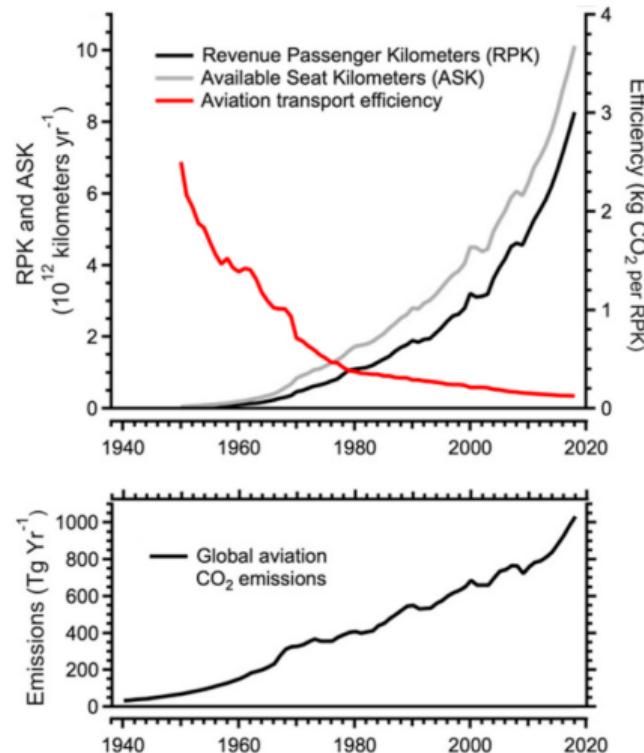
# Is Improved Efficiency Always Beneficial?

## Jevons Paradox (William S. Jevons, 1865)

Increases in efficiency with which a resource is used tend to increase the rate of consumption of that resource due to increasing demand.

### Examples:

1. Coal use (Jevons' original observation)
2. LED lights
3. Air travel [1]



[1] D. S. Lee *et al.*, "The contribution of global aviation to anthropogenic climate forcing for 2000 to 2018," 2021.

# Is Improved Efficiency Always Beneficial?

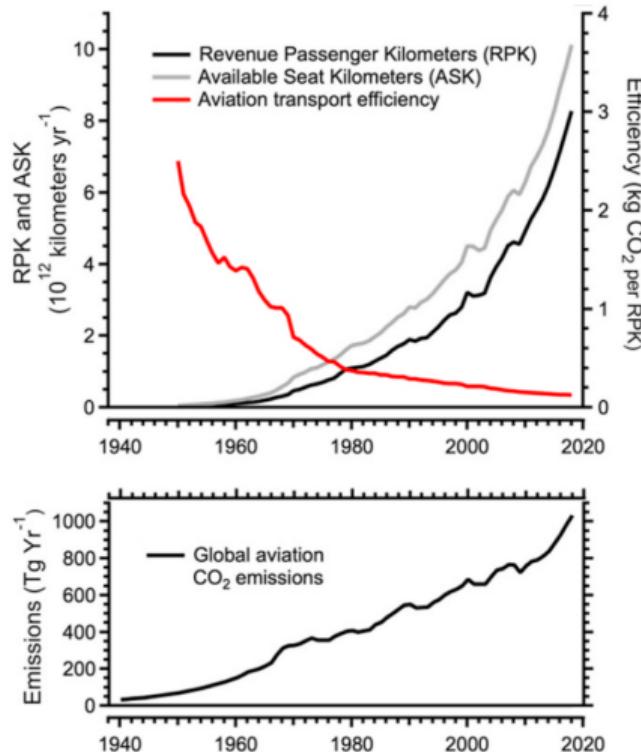
## Jevons Paradox (William S. Jevons, 1865)

Increases in efficiency with which a resource is used tend to increase the rate of consumption of that resource due to increasing demand.

### Examples:

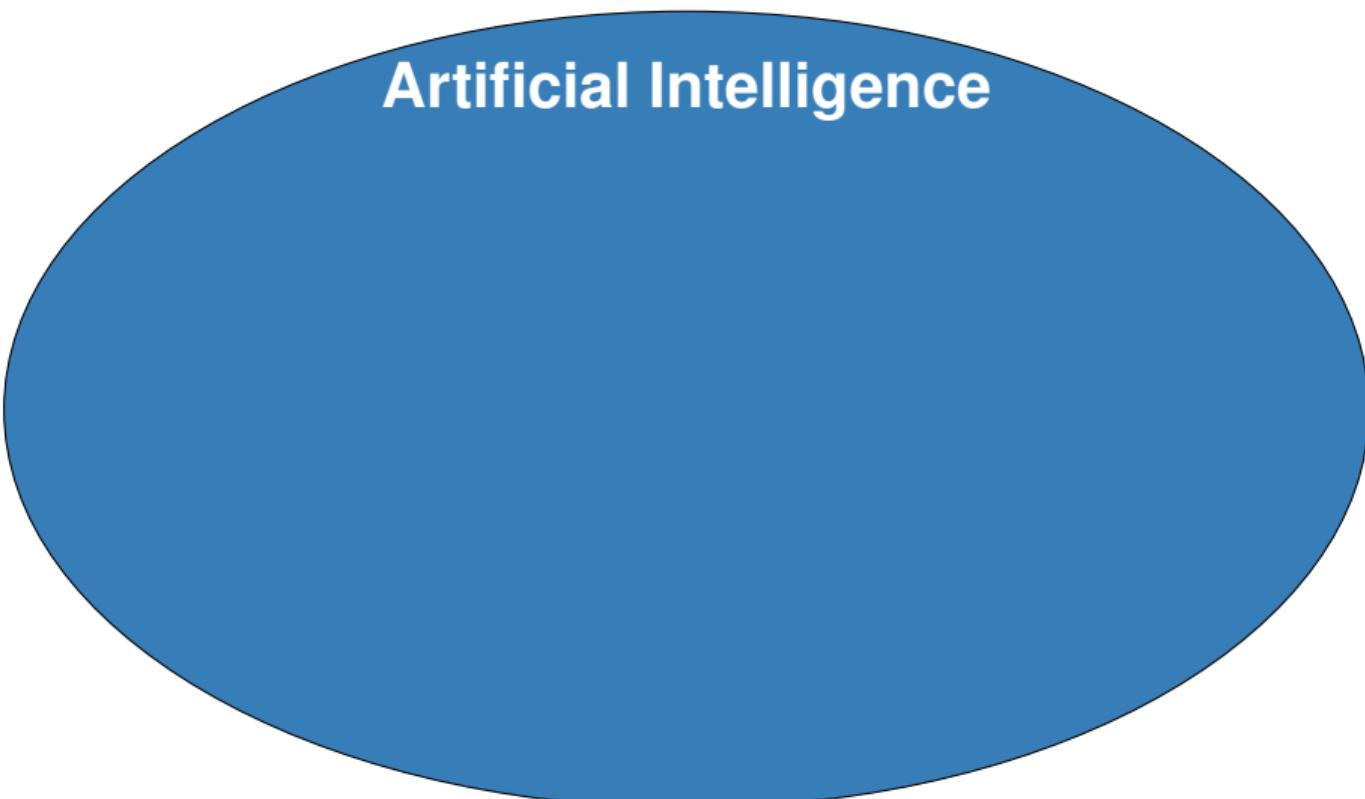
1. Coal use (Jevons' original observation)
2. LED lights
3. Air travel [1]

**So:** Don't look at energy efficiency in isolation!



[1] D. S. Lee *et al.*, "The contribution of global aviation to anthropogenic climate forcing for 2000 to 2018," 2021.

“...deep learning algorithms...”



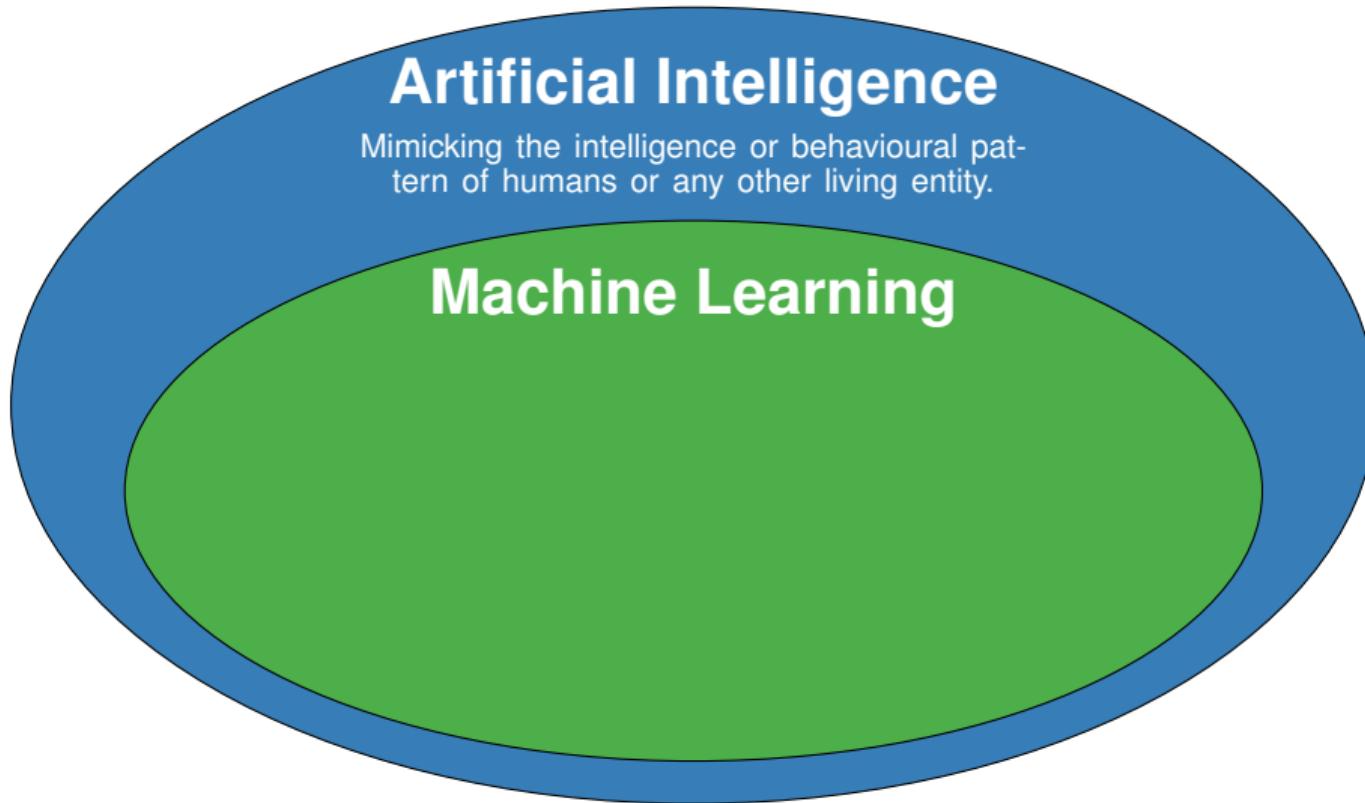
Artificial Intelligence

“...deep learning algorithms...”

## Artificial Intelligence

Mimicking the intelligence or behavioural pattern of humans or any other living entity.

“...deep learning algorithms...”



**“...deep learning algorithms...”**

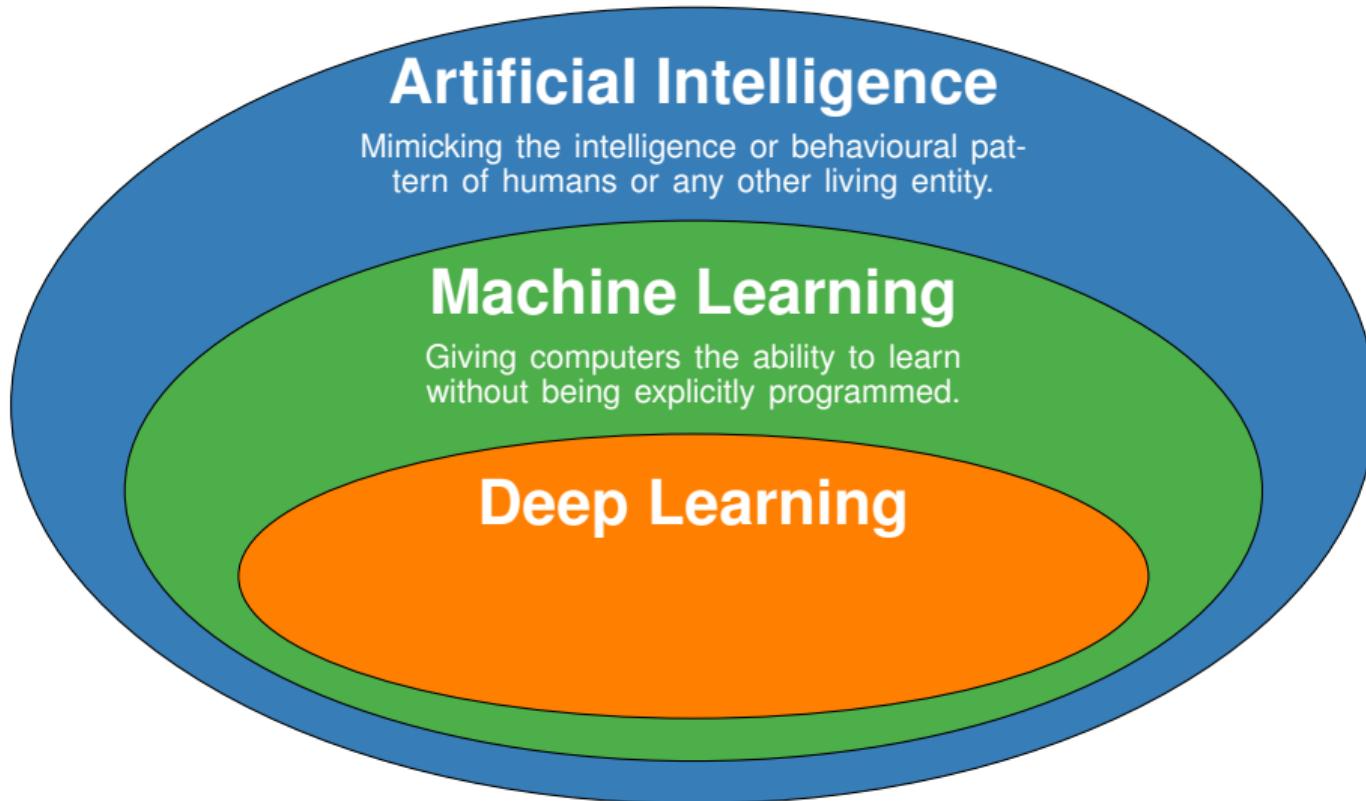
## **Artificial Intelligence**

Mimicking the intelligence or behavioural pattern of humans or any other living entity.

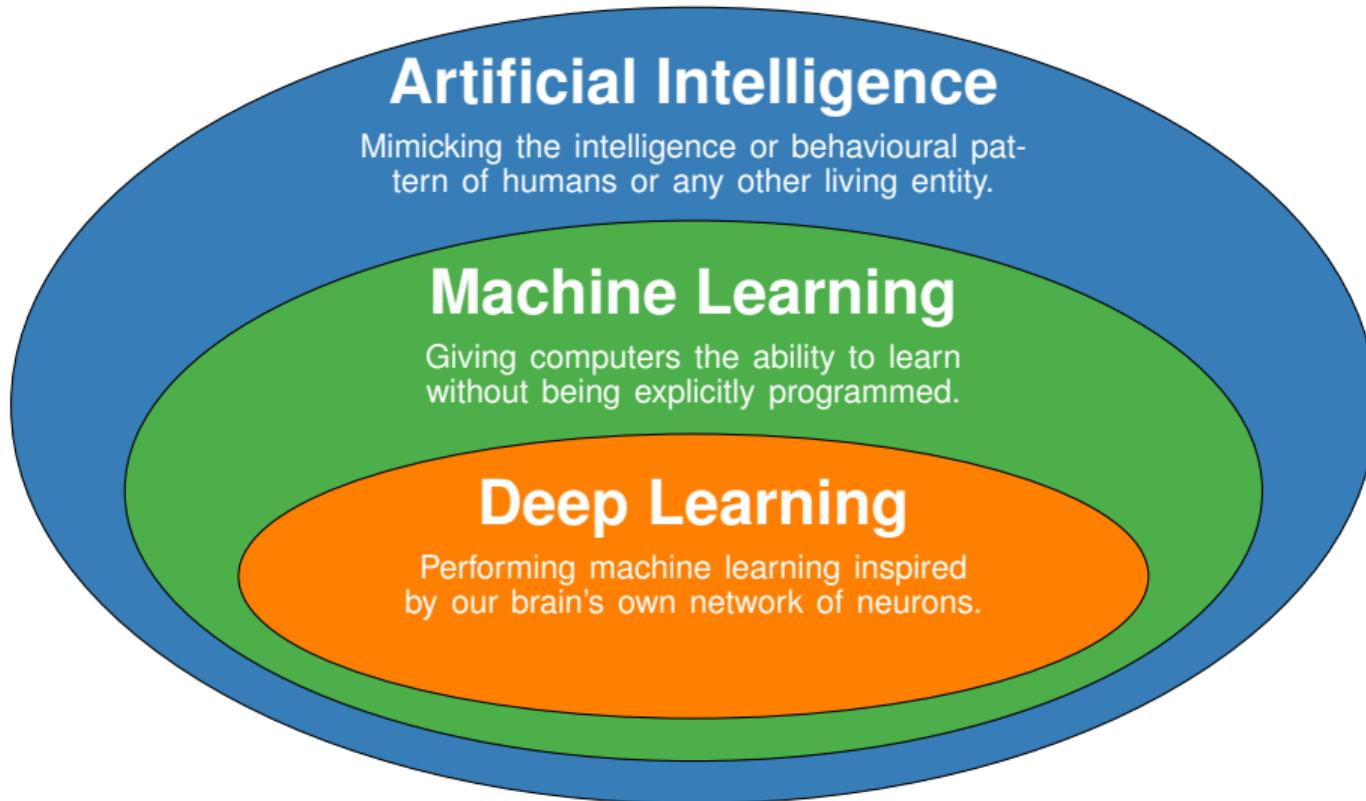
## **Machine Learning**

Giving computers the ability to learn without being explicitly programmed.

**“...deep learning algorithms...”**



**“...deep learning algorithms...”**



# Types of Learning

1. **Supervised learning:** learns from labeled datasets.
  - **Example:** Camera-based object detection for self-driving cars.

# Types of Learning

1. **Supervised learning:** learns from labeled datasets.
  - **Example:** Camera-based object detection for self-driving cars.
  
2. **Unsupervised learning:** discovers patterns in unlabeled datasets.
  - **Example:** Network intrusion detection via anomalies.

# Types of Learning

1. **Supervised learning:** learns from labeled datasets.
  - **Example:** Camera-based object detection for self-driving cars.
2. **Unsupervised learning:** discovers patterns in unlabeled datasets.
  - **Example:** Network intrusion detection via anomalies.
3. **Reinforcement learning:** learns to act based on feedback.
  - **Example:** Learning how to play Starcraft.

# Types of Learning

1. **Supervised learning:** learns from labeled datasets.
  - **Example:** Camera-based object detection for self-driving cars.
2. **Unsupervised learning:** discovers patterns in unlabeled datasets.
  - **Example:** Network intrusion detection via anomalies.
3. **Reinforcement learning:** learns to act based on feedback.
  - **Example:** Learning how to play Starcraft.

In practice, a combination of approaches is often used.

For example, ChatGPT uses both supervised learning and reinforcement learning.

# Learning Systems

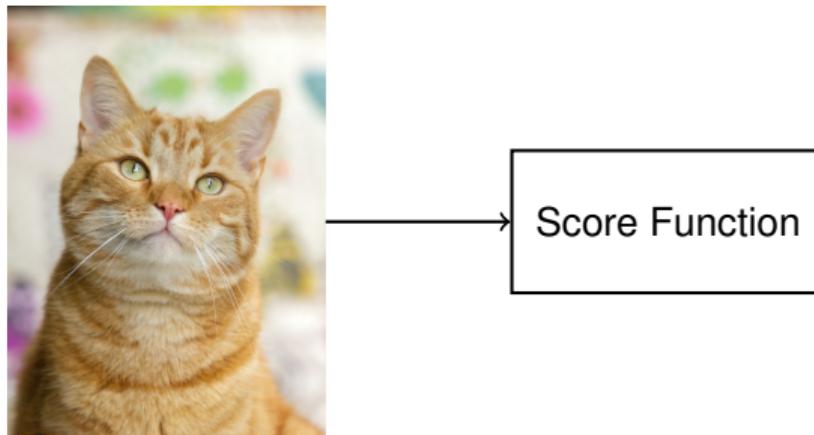
Basic components of a (supervised) learning system:



# Learning Systems

Basic components of a (supervised) learning system:

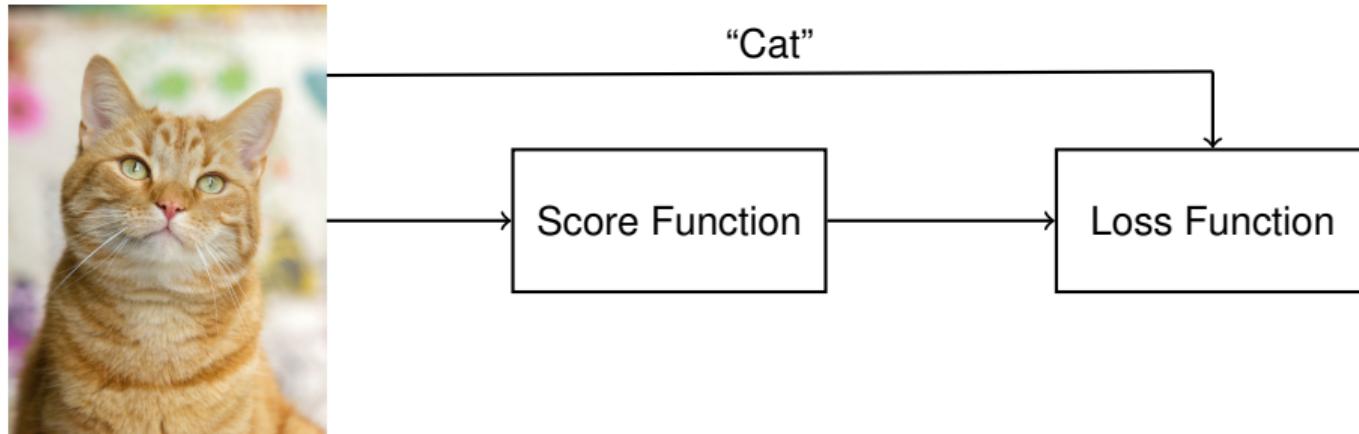
1. **Score function**



# Learning Systems

Basic components of a (supervised) learning system:

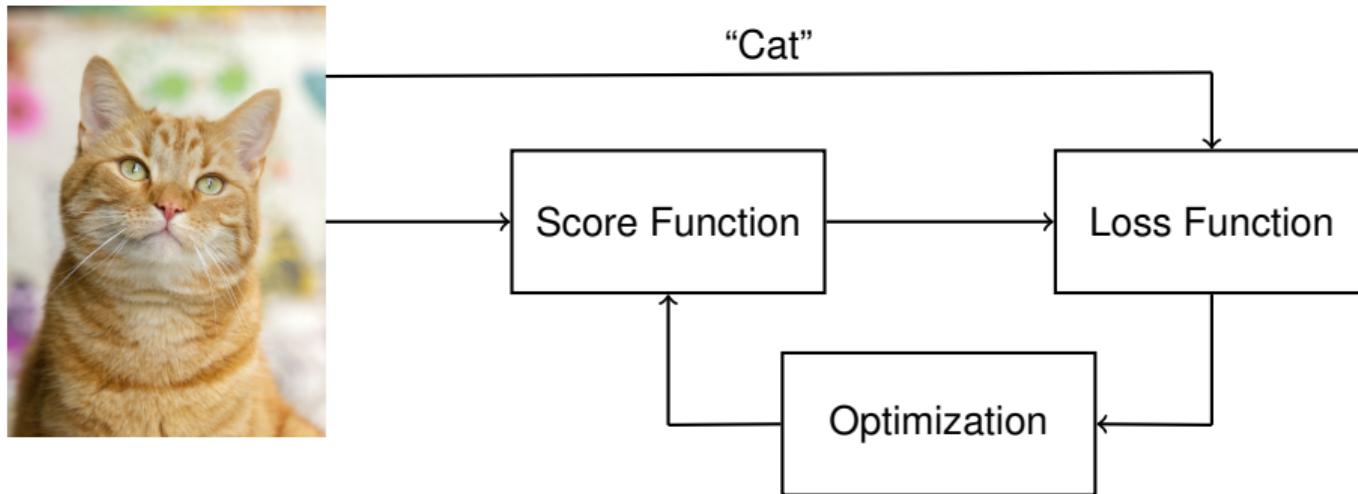
1. **Score function**
2. **Loss function**



# Learning Systems

Basic components of a (supervised) learning system:

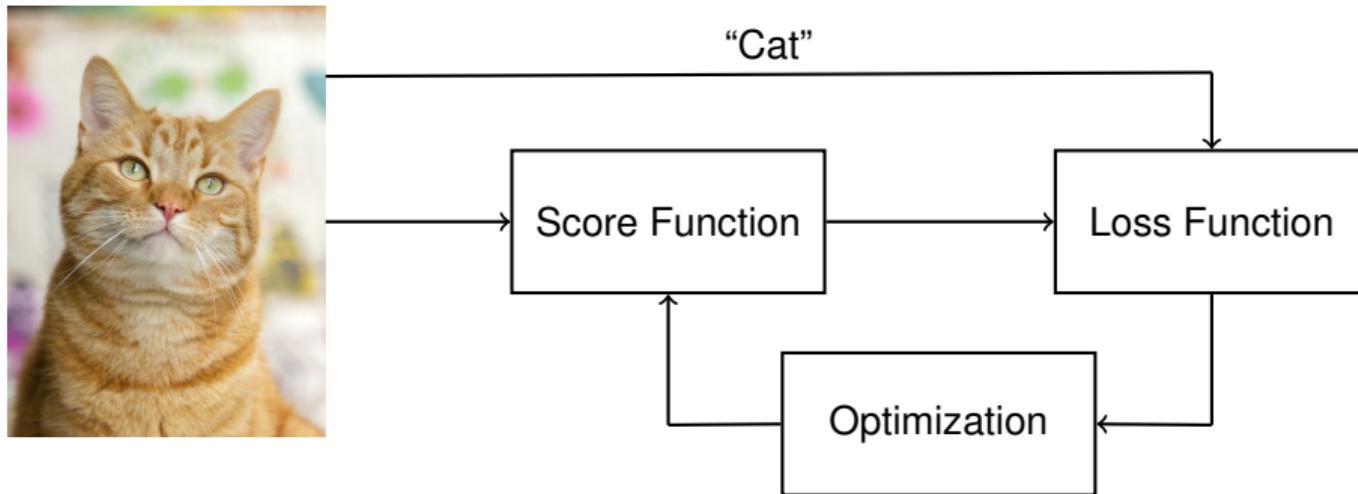
1. **Score function**
2. **Loss function**
3. **Optimization**



# Learning Systems

Basic components of a (supervised) learning system:

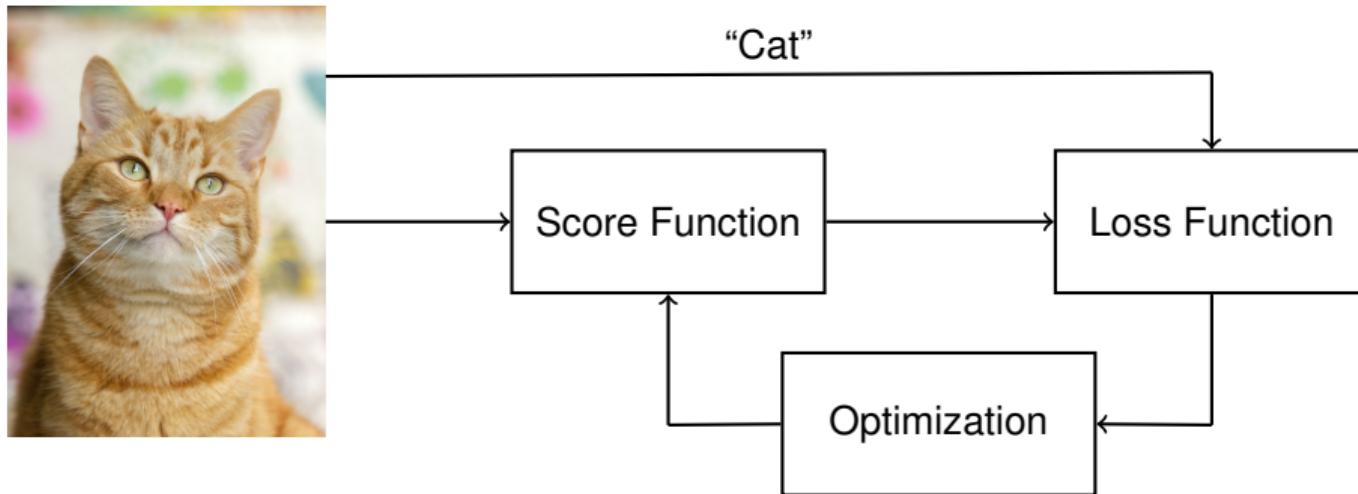
1. **Score function:** decision trees, support vector machines, neural networks, ...
2. **Loss function**
3. **Optimization**



# Learning Systems

Basic components of a (supervised) learning system:

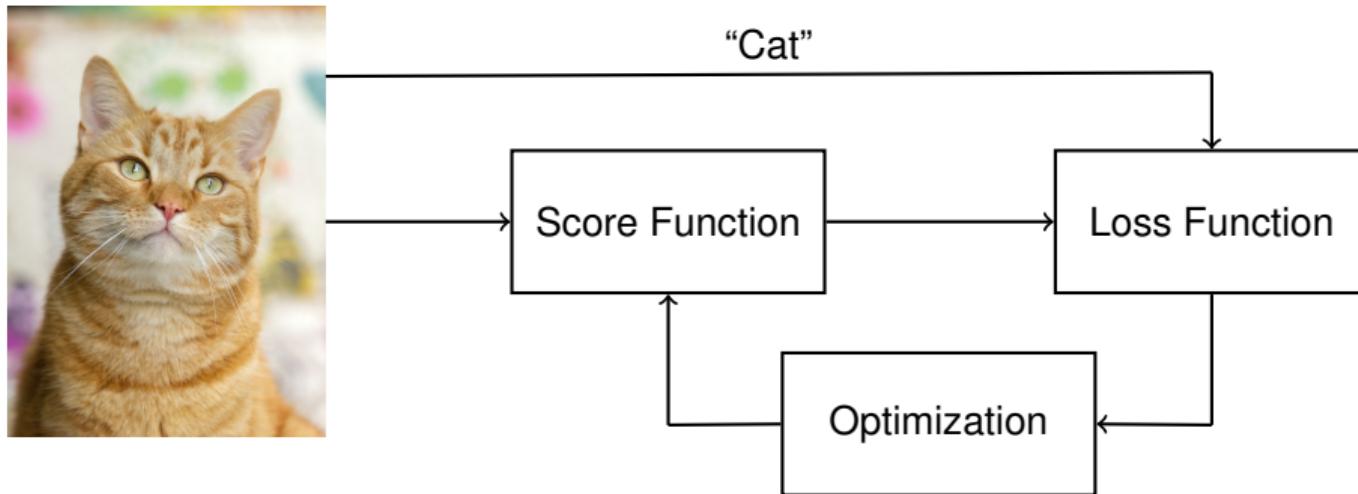
1. **Score function:** decision trees, support vector machines, **neural networks**, ...
2. **Loss function**
3. **Optimization**



# Learning Systems

Basic components of a (supervised) learning system:

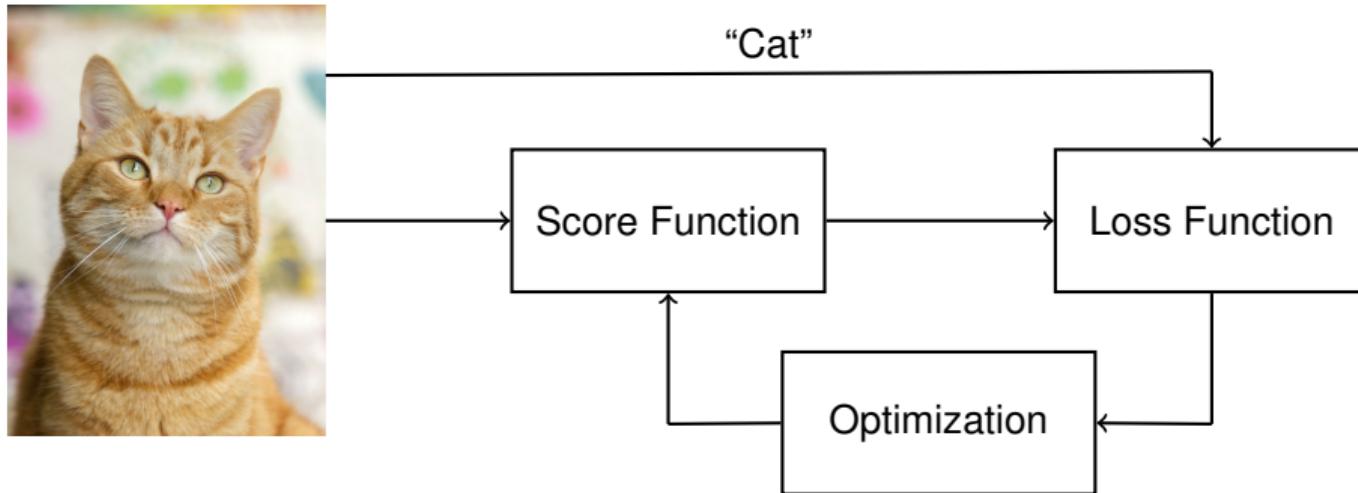
1. **Score function:** decision trees, support vector machines, **neural networks**, ...
2. **Loss function:** mean squared-error, binary cross-entropy, ...
3. **Optimization**



# Learning Systems

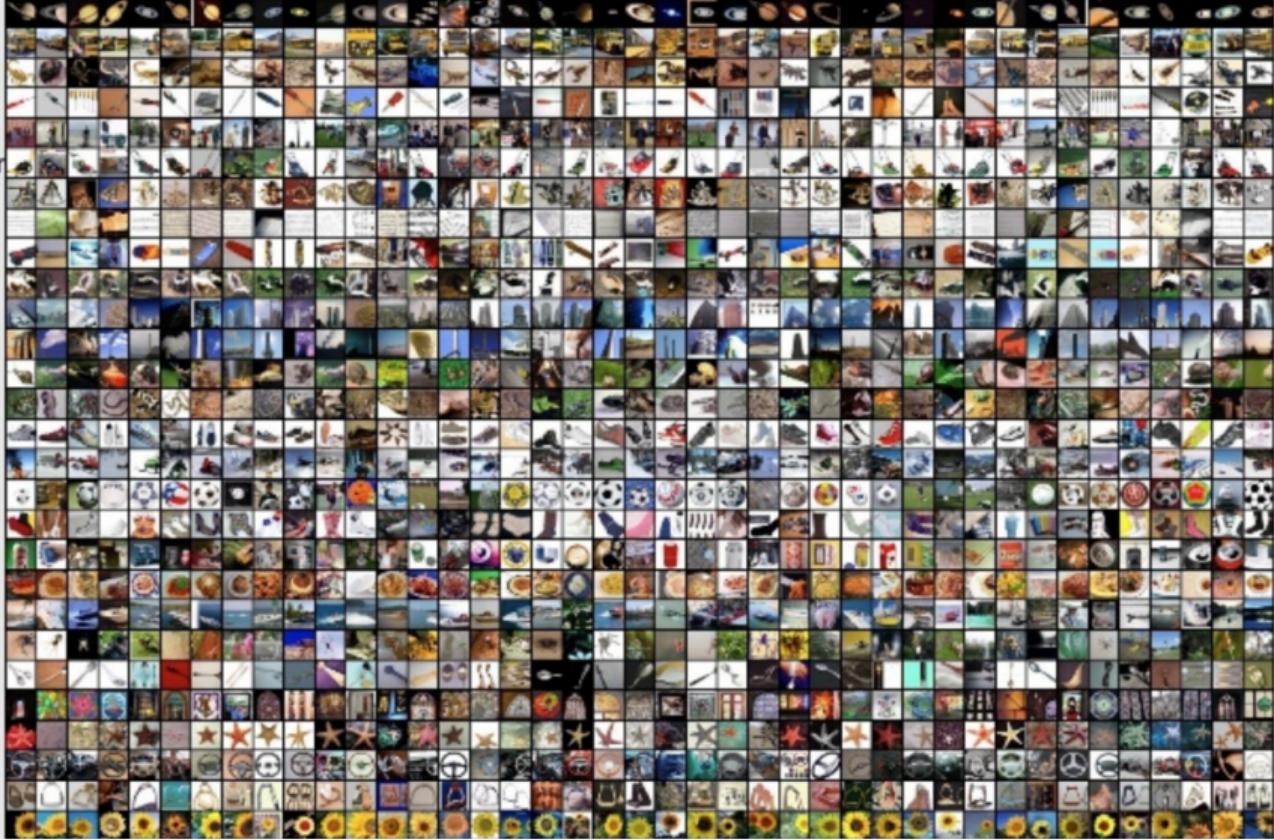
Basic components of a (supervised) learning system:

1. **Score function:** decision trees, support vector machines, **neural networks**, ...
2. **Loss function:** mean squared-error, binary cross-entropy, ...
3. **Optimization:** stochastic gradient descent, batch gradient descent, ...

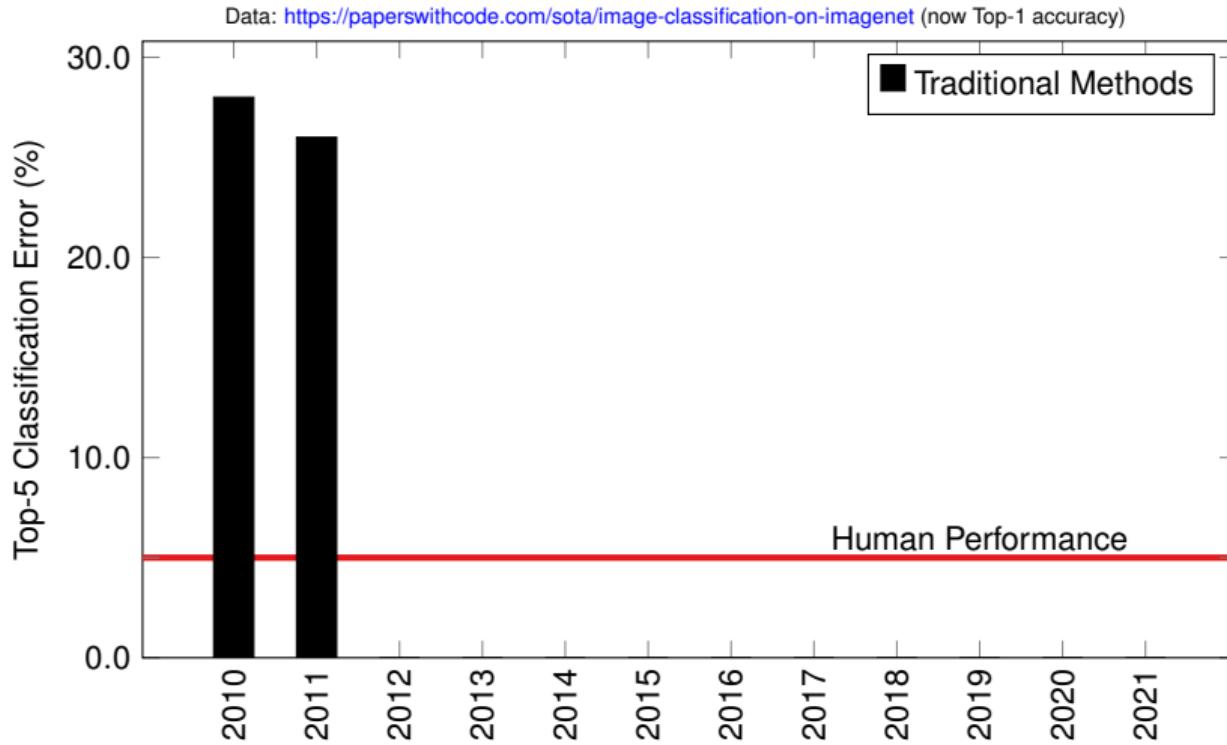


# The ImageNet Challenge

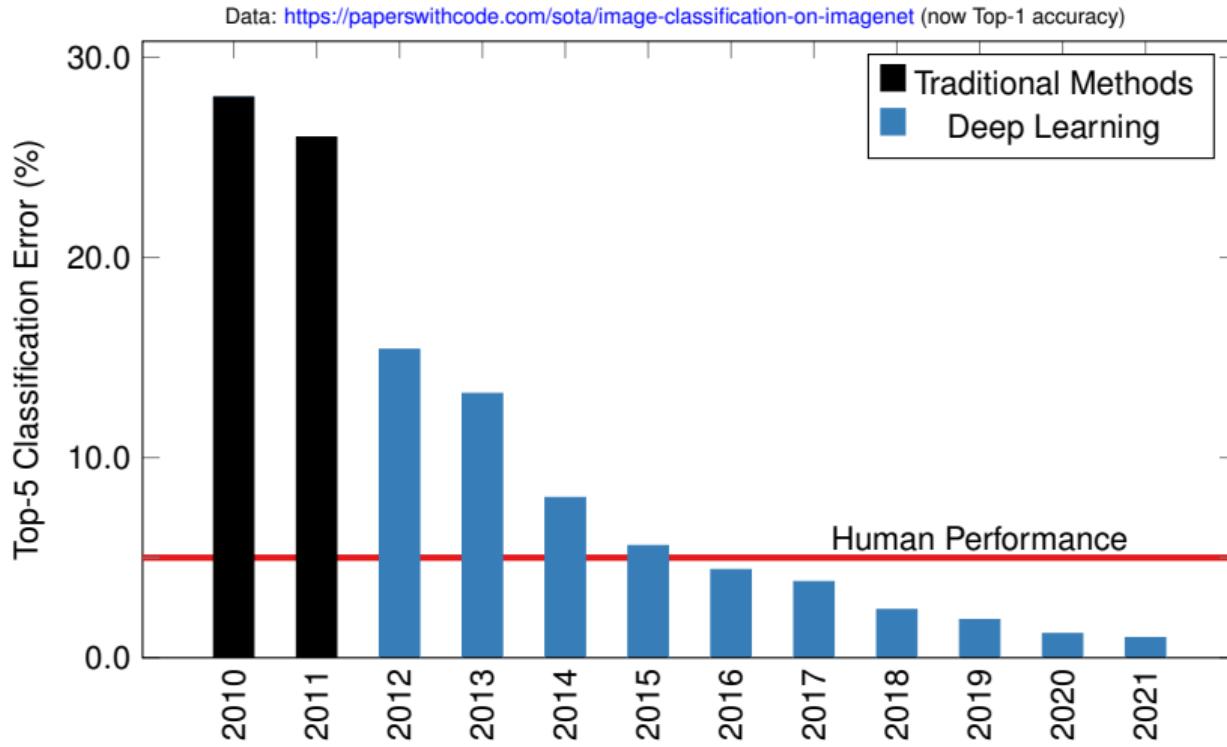
saturn  
school-bus  
scorpion-101  
screwdriver  
segway  
self-propelled-lawnmower  
sextant  
sheet-music  
skateboard  
skunk  
skyscraper  
smokestack  
snail  
snake  
sneaker  
snowmobile  
soccer-ball  
socks  
soda-can  
spaghetti  
speed-boat  
spider  
spoon  
stained-glass  
starfish-101  
steering-wheel  
stirrups  
sunflower-101



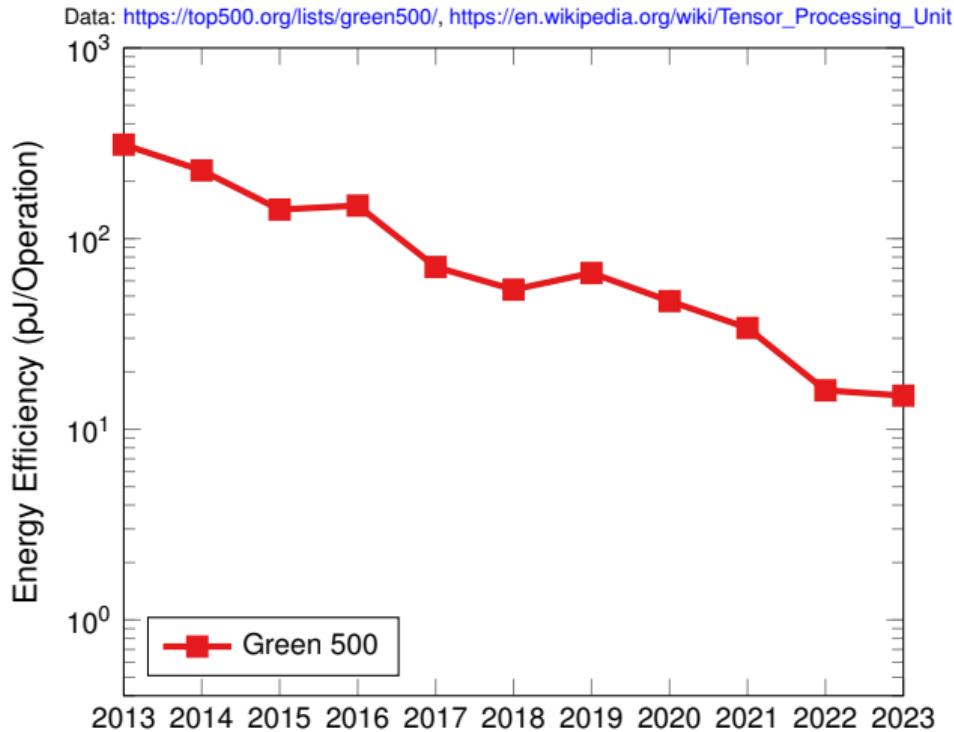
# The ImageNet Challenge



# The ImageNet Challenge



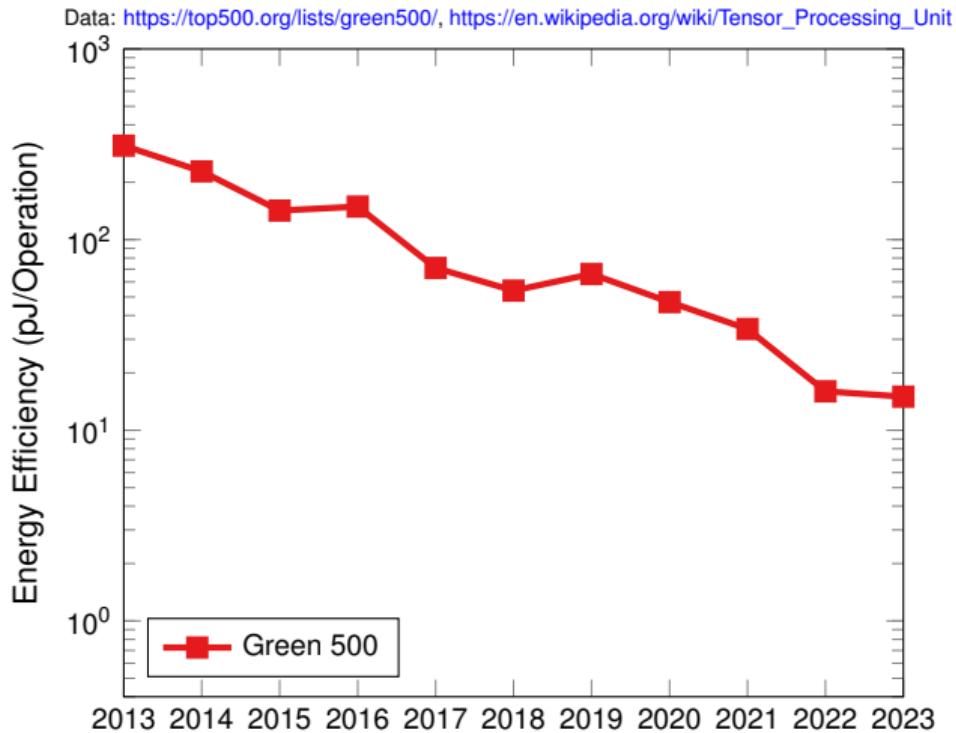
# “...and hardware architectures.”



The Frontier supercomputer.

OLCF at ORNL, CC BY 2.0, Wikimedia

# “...and hardware architectures.”

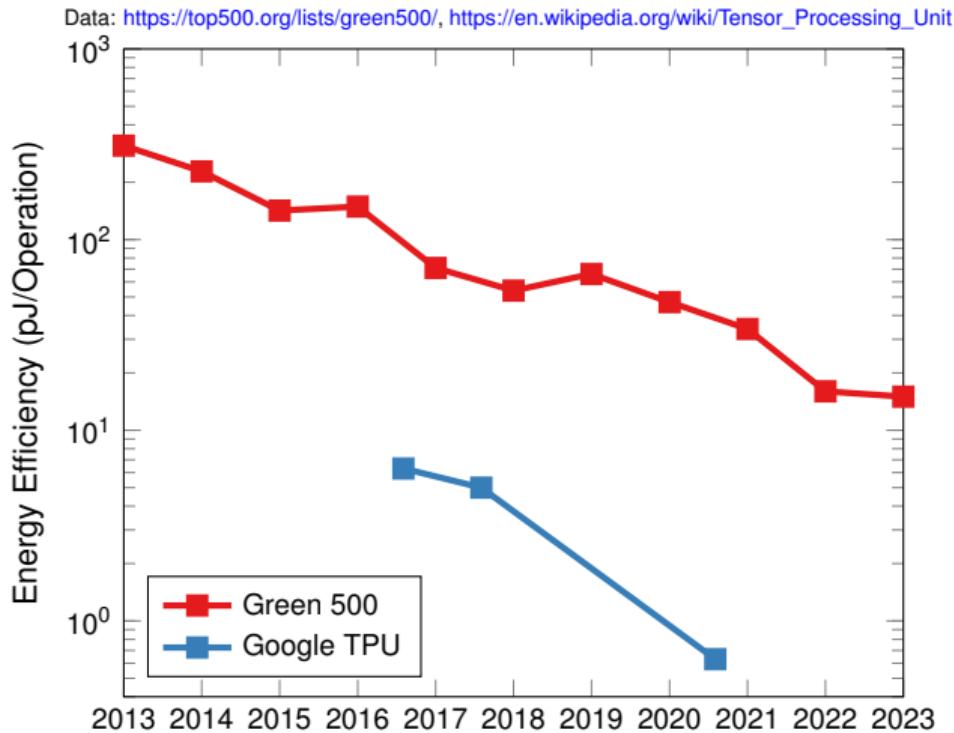


**Energy efficiency of brain:  $10^{-8}$  pJ/Operation!!!**



OLCF at ORNL, CC BY 2.0, Wikimedia

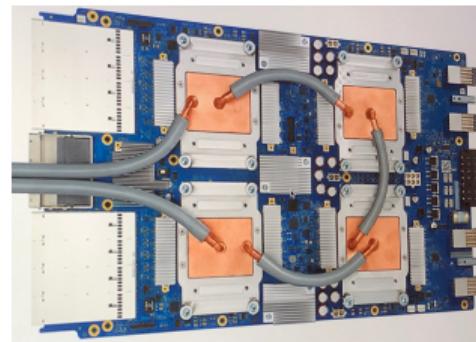
# “...and hardware architectures.”



**Energy efficiency of brain:  $10^{-8}$  pJ/Operation!!!**



The Frontier supercomputer.

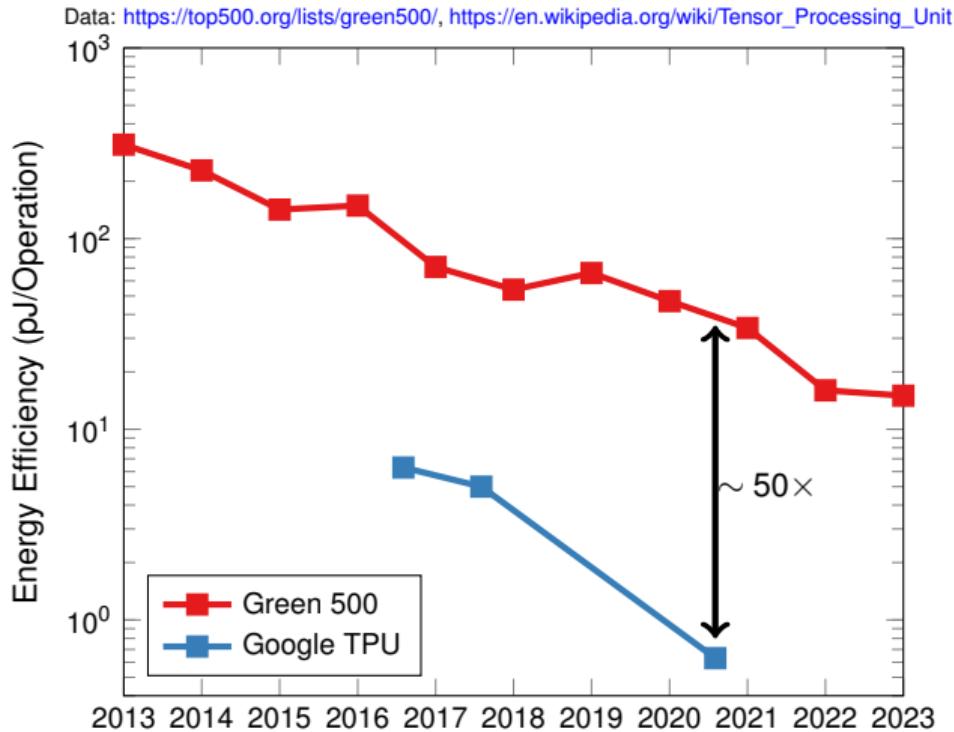


A Tensor Processing Unit.

OLCF at ORNL, CC BY 2.0, Wikimedia

Zinskauf, CC BY-SA 4.0, Wikimedia

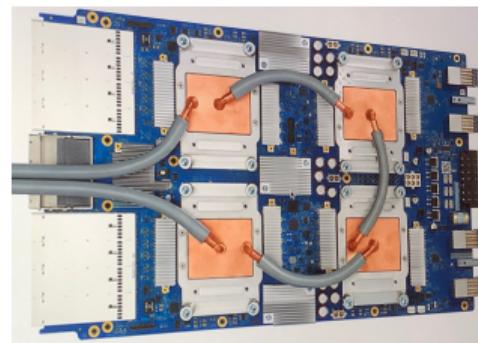
# “...and hardware architectures.”



**Energy efficiency of brain:  $10^{-8}$  pJ/Operation!!!**



The Frontier supercomputer.



A Tensor Processing Unit.

OLCF at ORNL, CC BY 2.0, Wikimedia

Zinskauf, CC BY-SA 4.0, Wikimedia

# Course Team

## Core Team



dr Alexios Balatsoukas-Stimming  
**Responsible lecturer**



dr Federico Corradi  
**Co-lecturer**



Zhanbo Shen  
**Teaching assistant**



Lorenzo Pes  
**Teaching assistant**

## Guest Lecturers



Prof. Henk Corporaal  
TU/e Emeritus



dr Maurice Peemen  
Thermo Fischer Scientific



dr Willem Sandberg  
NXP

# Learning Objectives & Prerequisites

At the end of this course, you will be able to:

1. Implement deep learning methods, including inference and learning, both from scratch and using modern frameworks.

# Learning Objectives & Prerequisites

At the end of this course, you will be able to:

1. Implement deep learning methods, including inference and learning, both from scratch and using modern frameworks.
2. Design deep neural network (DNN) architectures, with a focus on state-of-the-art convolutional neural networks (CNNs).

# Learning Objectives & Prerequisites

At the end of this course, you will be able to:

1. Implement deep learning methods, including inference and learning, both from scratch and using modern frameworks.
2. Design deep neural network (DNN) architectures, with a focus on state-of-the-art convolutional neural networks (CNNs).
3. Optimize DNNs using various methods (e.g., quantization, pruning) and explain the trade-offs of these methods.

# Learning Objectives & Prerequisites

At the end of this course, you will be able to:

1. Implement deep learning methods, including inference and learning, both from scratch and using modern frameworks.
2. Design deep neural network (DNN) architectures, with a focus on state-of-the-art convolutional neural networks (CNNs).
3. Optimize DNNs using various methods (e.g., quantization, pruning) and explain the trade-offs of these methods.
4. Evaluate special processing architectures and dedicated hardware for efficient deep learning.

# Learning Objectives & Prerequisites

At the end of this course, you will be able to:

1. Implement deep learning methods, including inference and learning, both from scratch and using modern frameworks.
2. Design deep neural network (DNN) architectures, with a focus on state-of-the-art convolutional neural networks (CNNs).
3. Optimize DNNs using various methods (e.g., quantization, pruning) and explain the trade-offs of these methods.
4. Evaluate special processing architectures and dedicated hardware for efficient deep learning.

## Prerequisites:

1. Basic linear algebra.
2. Basic differential calculus.
3. Very good programming skills (ideally in Python).

# Course Organization

- **Lectures:** Tuesdays 10:45-12:30 & Fridays 15:30-17:15
- **TA hours:** Tuesdays 12:30-13:30 (Flux 4.076) & Thursdays 12:30-13:30 (Flux 4.078)

# Course Organization

- **Lectures:** Tuesdays 10:45-12:30 & Fridays 15:30-17:15
- **TA hours:** Tuesdays 12:30-13:30 (Flux 4.076) & Thursdays 12:30-13:30 (Flux 4.078)
- **Assessment:**
  - 70% written final exam (includes material from lectures & assignments)
  - 30% assignments (group code submission and **individual single-attempt** quizzes)
  - Minimum of 5 in each component.

# Course Organization

- **Lectures:** Tuesdays 10:45-12:30 & Fridays 15:30-17:15
- **TA hours:** Tuesdays 12:30-13:30 (Flux 4.076) & Thursdays 12:30-13:30 (Flux 4.078)
- **Assessment:**
  - 70% written final exam (includes material from lectures & assignments)
  - 30% assignments (group code submission and **individual single-attempt** quizzes)
  - Minimum of 5 in each component.
- **Assignments:**
  - **Assignment 1 (Basics):** Neural network inference and learning.
  - **Assignment 2 (Optimizations):** Quantization, pruning, neural architecture search.
  - **Assignment 3 (Platforms):** Hardware accelerators.

# Course Organization

- **Lectures:** Tuesdays 10:45-12:30 & Fridays 15:30-17:15
- **TA hours:** Tuesdays 12:30-13:30 (Flux 4.076) & Thursdays 12:30-13:30 (Flux 4.078)
- **Assessment:**
  - 70% written final exam (includes material from lectures & assignments)
  - 30% assignments (group code submission and **individual single-attempt** quizzes)
  - Minimum of 5 in each component.
- **Assignments:**
  - **Assignment 1 (Basics):** Neural network inference and learning.
  - **Assignment 2 (Optimizations):** Quantization, pruning, neural architecture search.
  - **Assignment 3 (Platforms):** Hardware accelerators.
- **Course material:** Canvas.

# Use of Generative AI

## Not allowed

Using AI to **create content** like images, source code, video, text or any other kind.

## Allowed

Use of AI-powered tools to improve **your own** content that operate at the level of at most a single sentence (e.g. spelling checkers) or a single line of code (as provided by many IDEs).

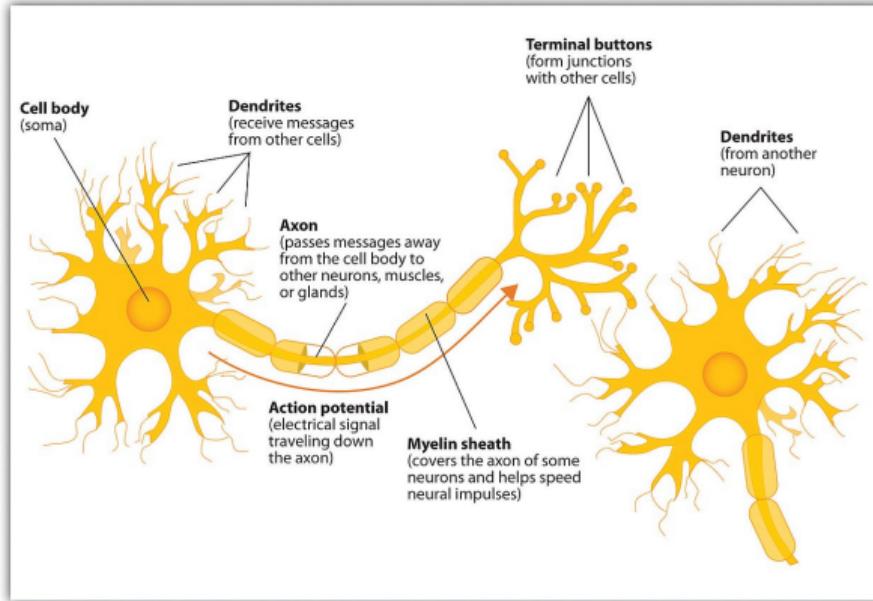
## Time For A Break

*“Our civilization is first and foremost a civilization of means; in the reality of modern life, the means, it would seem, are more important than the ends.”*

– Jacques Ellul, “[The Technological Society](#),” 1954.

# Our Brain

Jennifer Walinga, CC BY-SA 4.0, via Wikimedia Commons

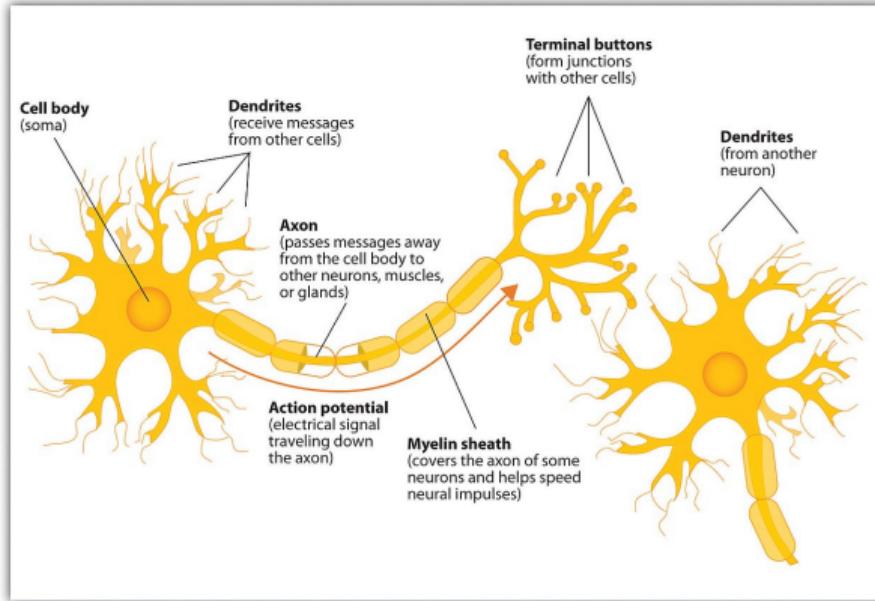


The basic computational unit of the brain is a **neuron**.

- Neurons receive input signals from **dendrites** and produce output signals along their **axon**, which interact with the dendrites of other neurons via **synaptic weights**.

# Our Brain

Jennifer Walinga, CC BY-SA 4.0, via Wikimedia Commons

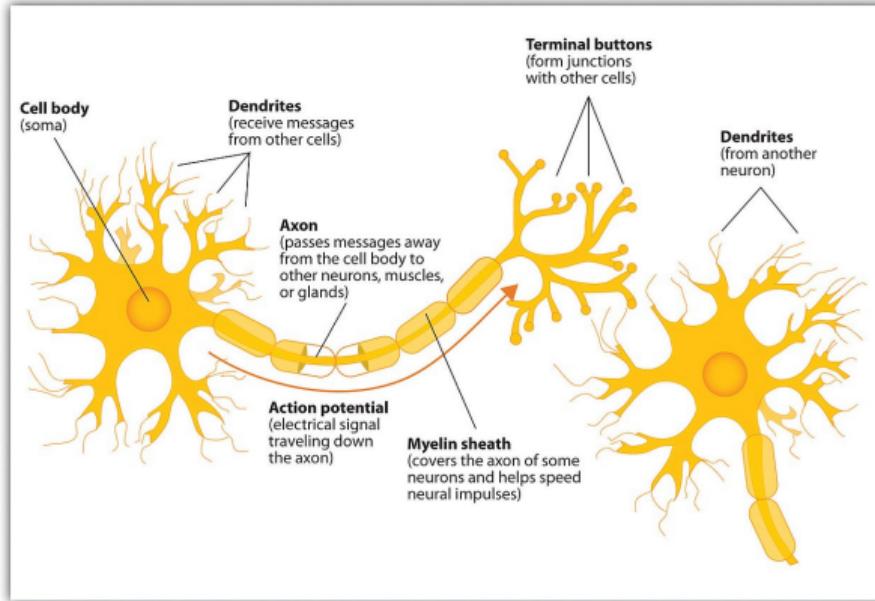


The basic computational unit of the brain is a **neuron**.

- Neurons receive input signals from **dendrites** and produce output signals along their **axon**, which interact with the dendrites of other neurons via **synaptic weights**.
- About 80 billion neurons.
- About  $10^{14}$  to  $10^{15}$  synapses.

# Our Brain

Jennifer Walinga, CC BY-SA 4.0, via Wikimedia Commons

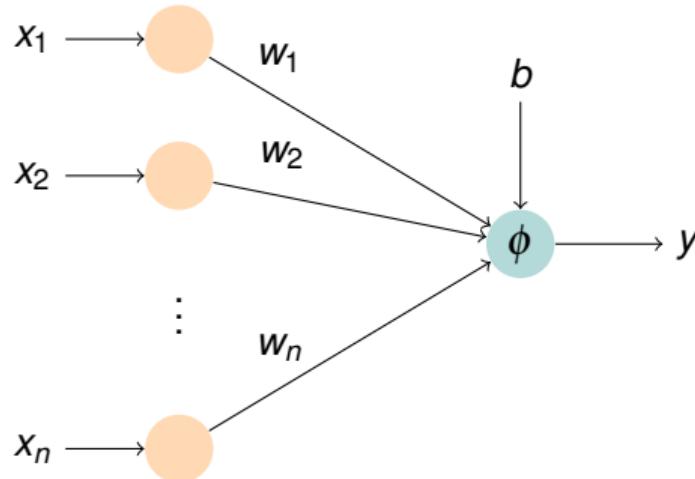


The basic computational unit of the brain is a **neuron**.

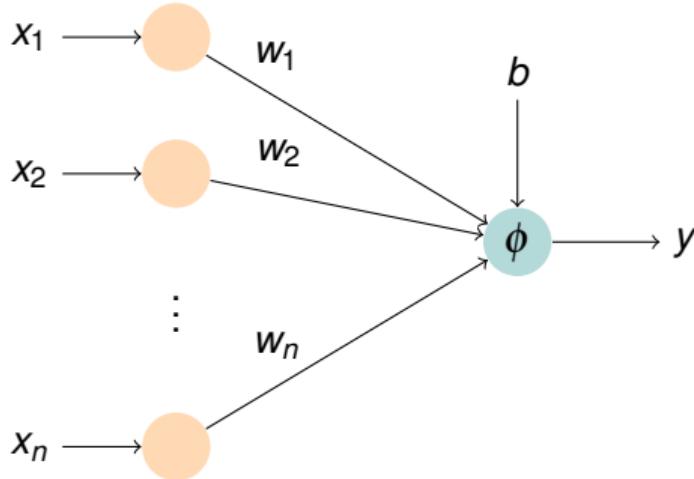
- Neurons receive input signals from **dendrites** and produce output signals along their **axon**, which interact with the dendrites of other neurons via **synaptic weights**.
- About 80 billion neurons.
- About  $10^{14}$  to  $10^{15}$  synapses.

The synaptic weights are **learnable** and they control the influence between neurons.

# An Artificial Neuron

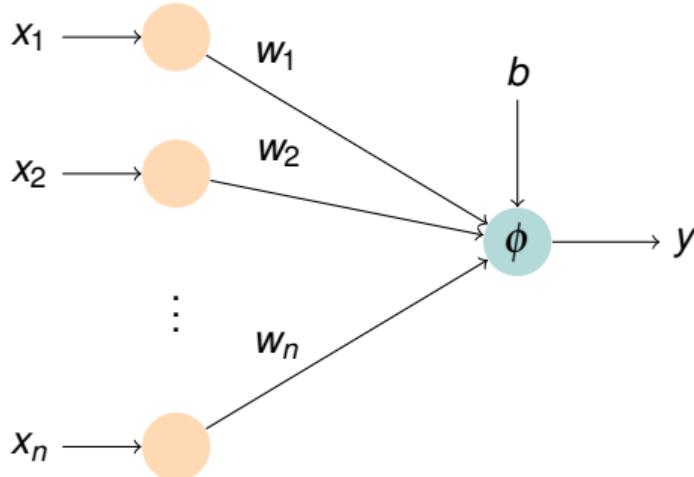


# An Artificial Neuron



$$y = \phi \left( \sum_{i=1}^n w_i x_i + b \right) = \phi (\mathbf{w}^\top \mathbf{x} + b)$$

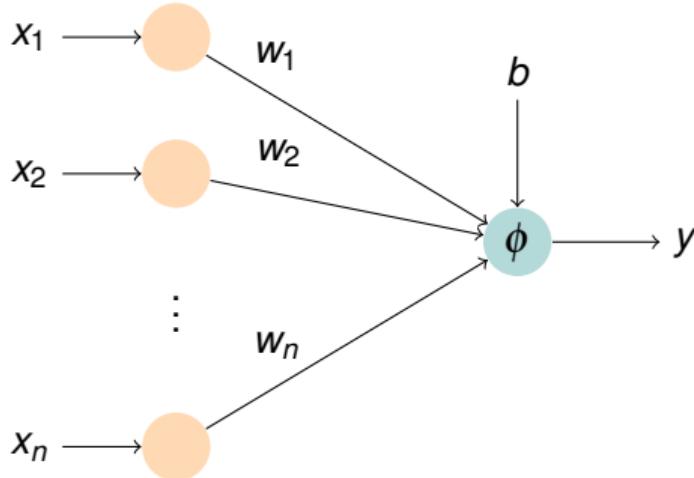
# An Artificial Neuron



- **Input** nodes correspond to dendrites.

$$y = \phi \left( \sum_{i=1}^n w_i x_i + b \right) = \phi (\mathbf{w}^\top \mathbf{x} + b)$$

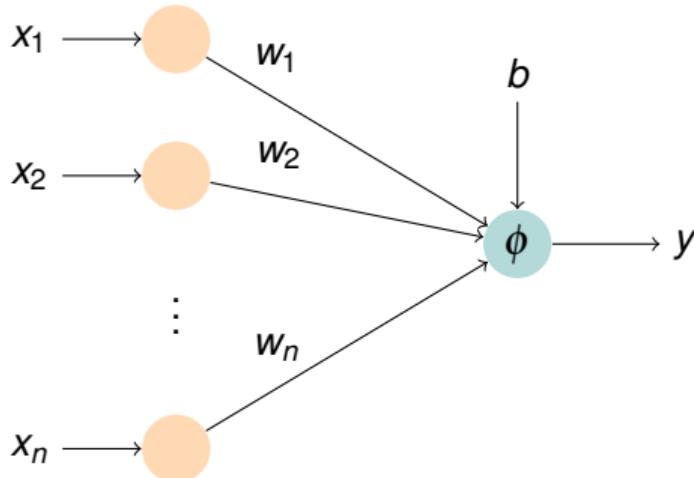
# An Artificial Neuron



- **Input** nodes correspond to dendrites.
- **Weights**  $w_i$  correspond to synaptic weights.

$$y = \phi \left( \sum_{i=1}^n w_i x_i + b \right) = \phi (\mathbf{w}^\top \mathbf{x} + b)$$

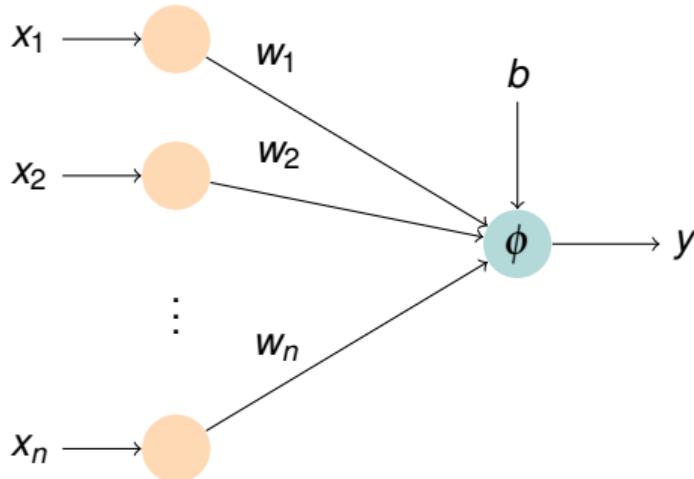
# An Artificial Neuron



- **Input** nodes correspond to dendrites.
- **Weights**  $w_i$  correspond to synaptic weights.
- **Hidden** node corresponds to soma-axon interaction:
  1. Incoming potentials and bias  $b$  are added up.
  2. A non-linear function  $\phi(\cdot)$  is applied.

$$y = \phi \left( \sum_{i=1}^n w_i x_i + b \right) = \phi (\mathbf{w}^\top \mathbf{x} + b)$$

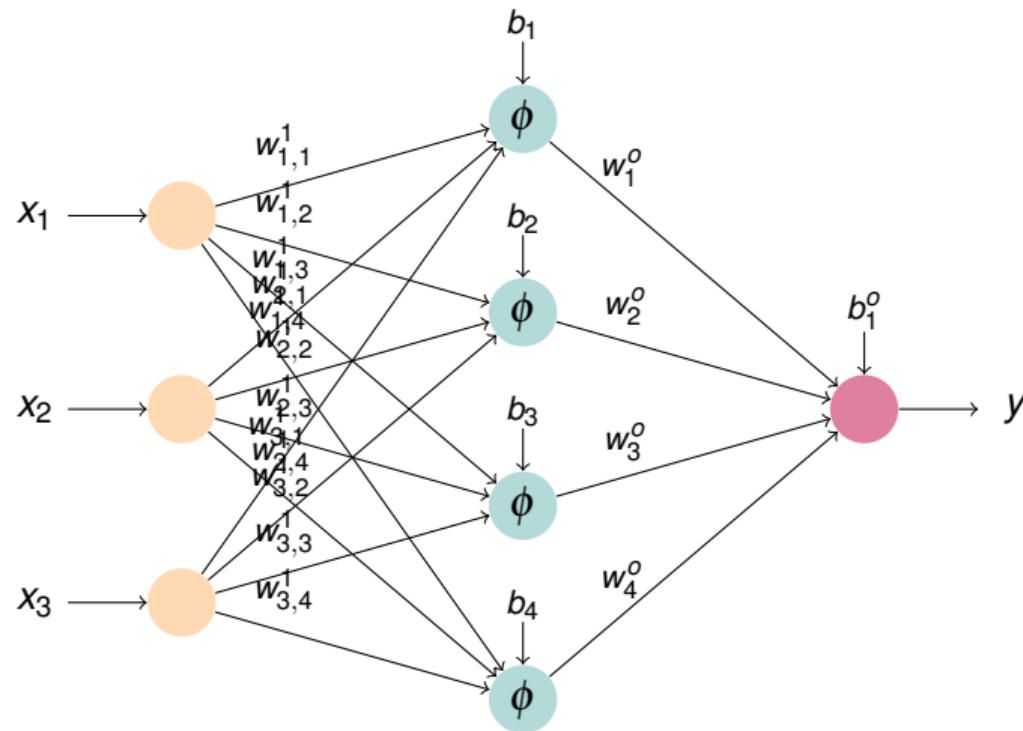
# An Artificial Neuron



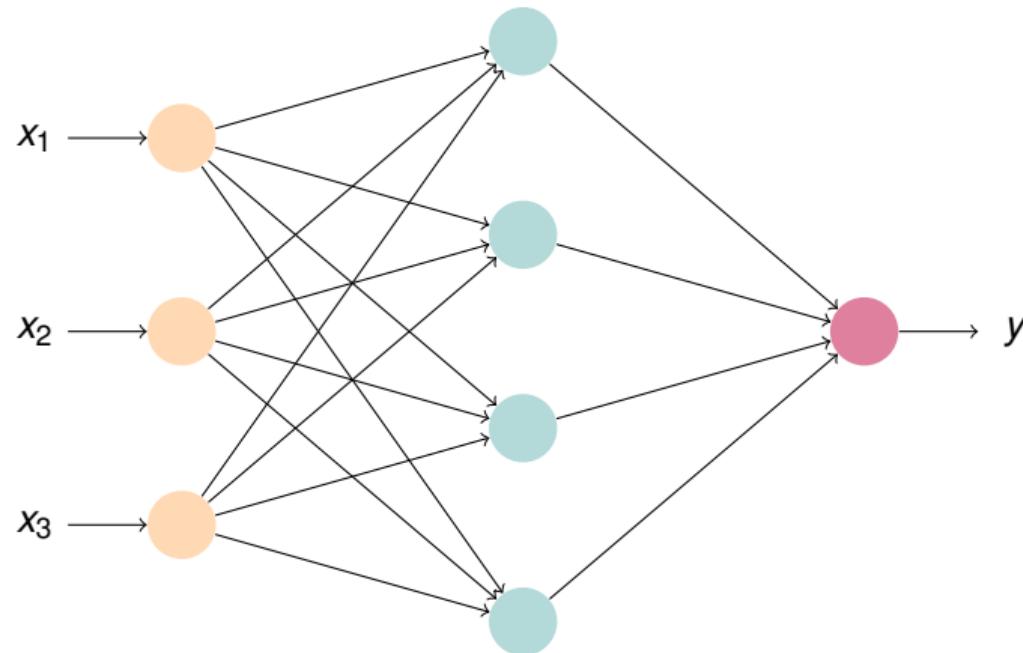
$$y = \phi \left( \sum_{i=1}^n w_i x_i + b \right) = \phi (\mathbf{w}^\top \mathbf{x} + b)$$

- **Input** nodes correspond to dendrites.
- **Weights**  $w_i$  correspond to synaptic weights.
- **Hidden** node corresponds to soma-axon interaction:
  1. Incoming potentials and bias  $b$  are added up.
  2. A non-linear function  $\phi(\cdot)$  is applied.
- **Output**  $y$  corresponds to action potential on axon.

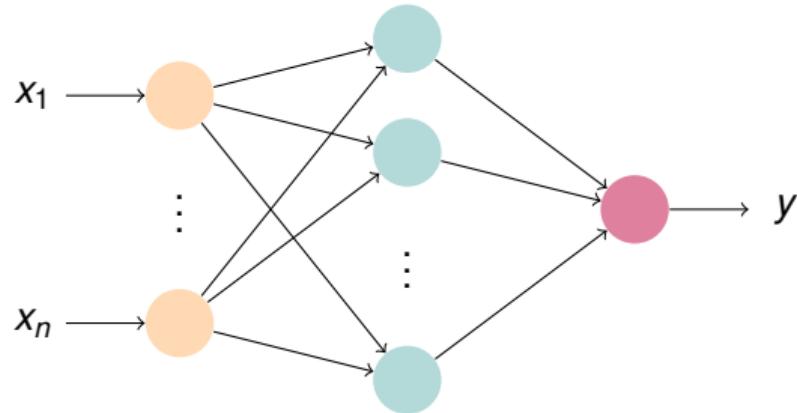
# A Single-Layer Single-Output Neural Network



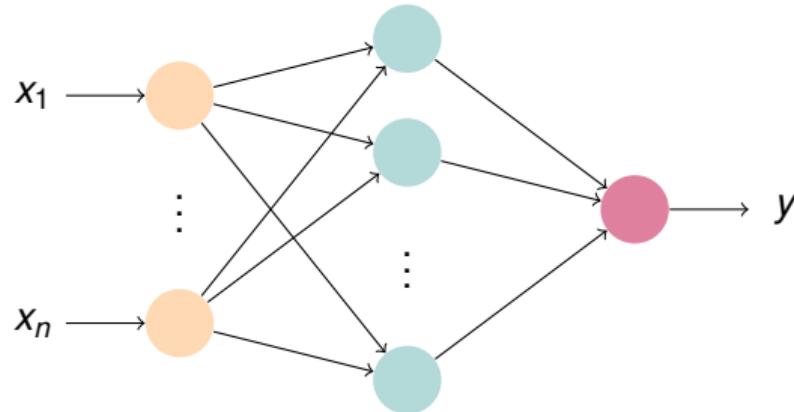
# A Single-Layer Single-Output Neural Network



# A Single-Layer Single-Output Neural Network

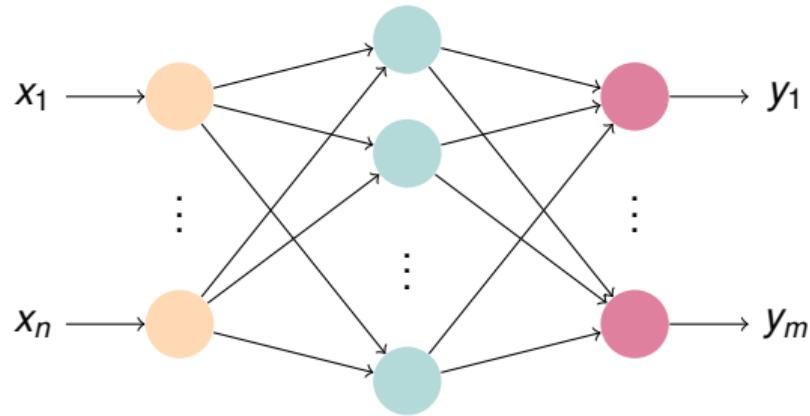


# A Single-Layer Single-Output Neural Network



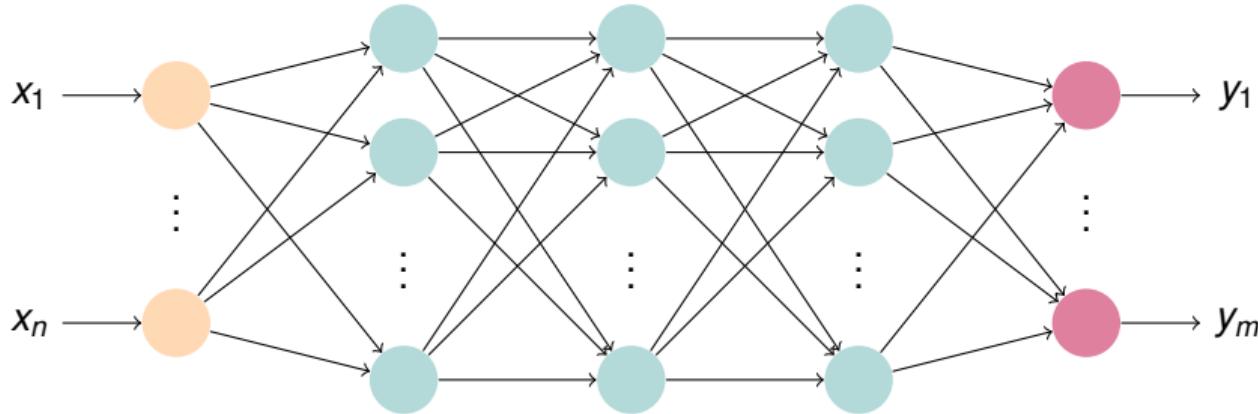
$$y = \left( \mathbf{w}^{[2]} \right)^\top \phi \left( \mathbf{W}^{[1]} \mathbf{x} + \mathbf{b}^{[1]} \right) + b^{[2]}$$

# A Single-Layer Multi-Output Neural Network



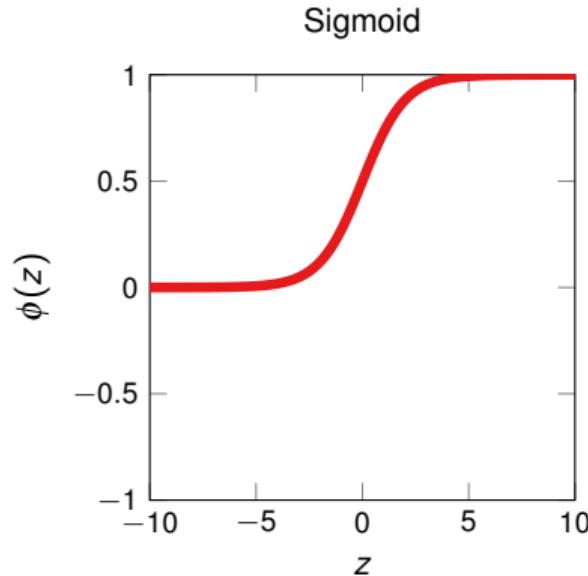
$$\mathbf{y} = \mathbf{W}^{[2]} \phi \left( \mathbf{W}^{[1]} \mathbf{x} + \mathbf{b}^{[1]} \right) + \mathbf{b}^{[2]}$$

# A Multi-Layer Multi-Output Neural Network



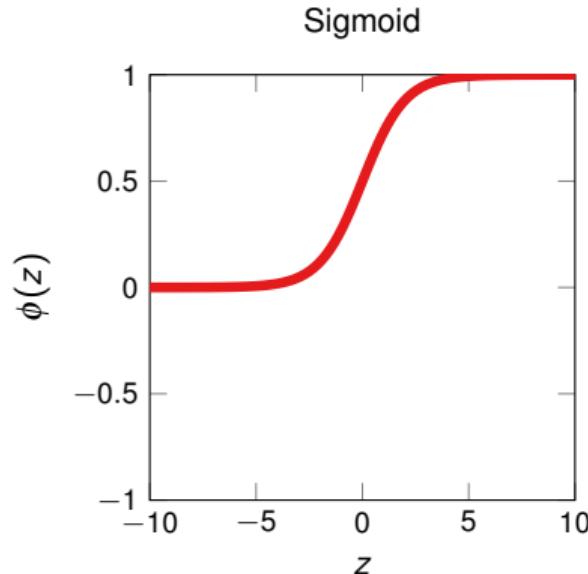
Let  $f^{[i]}(\mathbf{z}) = \phi \left( \mathbf{W}^{[i]}\mathbf{z} + \mathbf{b}^{[i]} \right)$ . Then, we have  $\mathbf{y} = f^{[4]} \left( f^{[3]} \left( f^{[2]} \left( f^{[1]}(\mathbf{x}) \right) \right) \right)$

# Traditional Activation Functions

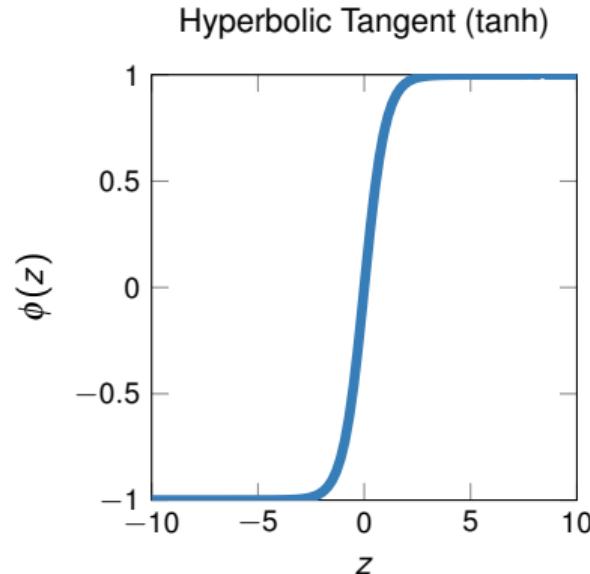


$$\phi(z) = \frac{1}{1+e^{-z}} = \sigma(z)$$

# Traditional Activation Functions



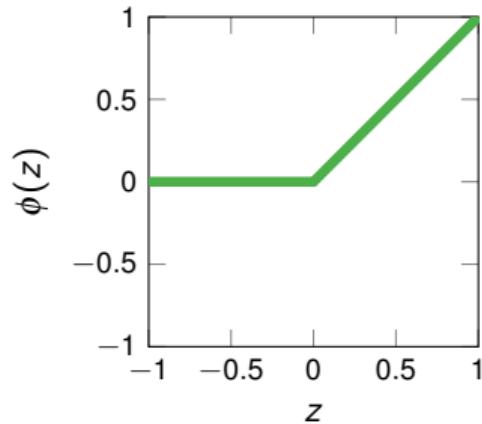
$$\phi(z) = \frac{1}{1+e^{-z}} = \sigma(z)$$



$$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

# Modern Activation Functions

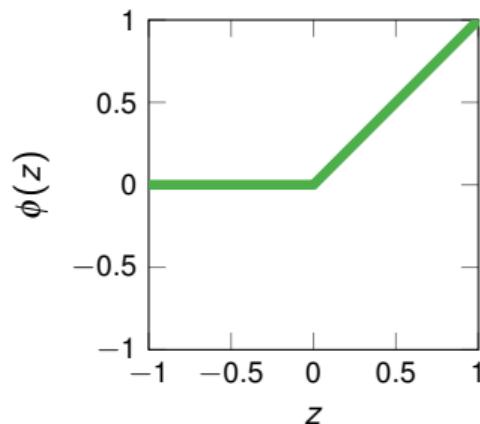
Rectified Linear Unit (ReLU)



$$\phi(z) = \max(0, z)$$

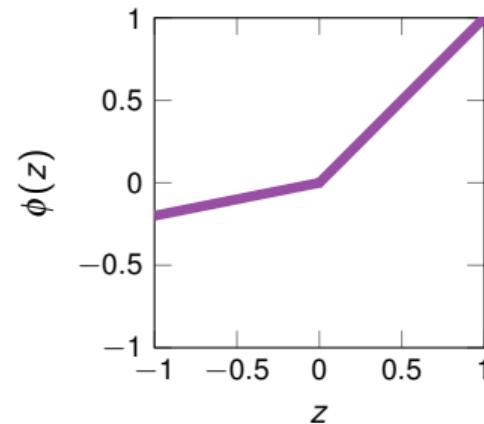
# Modern Activation Functions

Rectified Linear Unit (ReLU)



$$\phi(z) = \max(0, z)$$

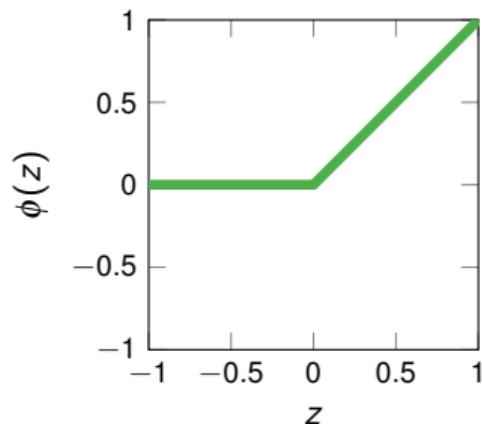
Leaky ReLU ( $\alpha = 0.2$ )



$$\phi(z) = \max(\alpha z, z)$$

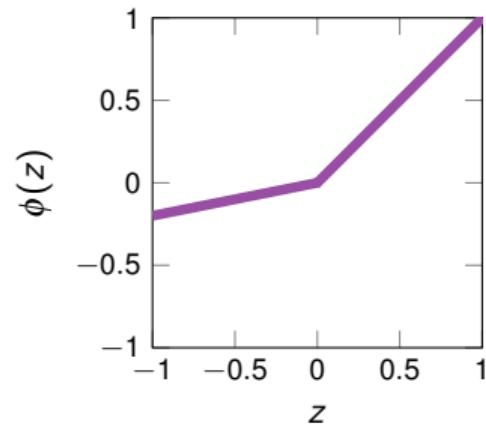
# Modern Activation Functions

Rectified Linear Unit (ReLU)



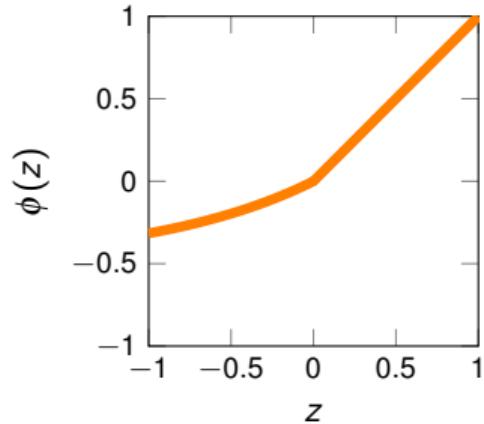
$$\phi(z) = \max(0, z)$$

Leaky ReLU ( $\alpha = 0.2$ )



$$\phi(z) = \max(\alpha z, z)$$

Exponential Linear Unit (ELU,  $\alpha = 0.5$ )



$$\phi(z) = \begin{cases} \alpha(e^z - 1), & z < 0, \\ z, & z \geq 0. \end{cases}$$

# Final Layer Activation Function

- The final layer activation function is chosen **depending on the task.**

# Final Layer Activation Function

- The final layer activation function is chosen **depending on the task**.
- For **regression**, typically a linear activation function (i.e.,  $\phi(z) = z$ ).

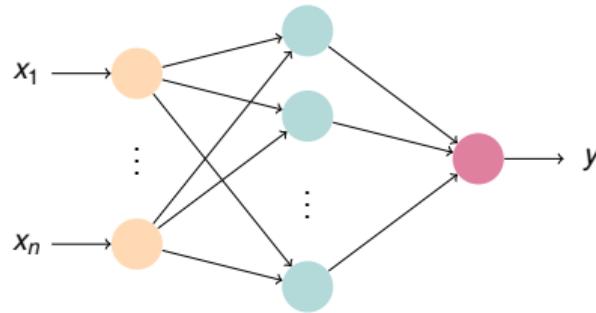
# Final Layer Activation Function

- The final layer activation function is chosen **depending on the task**.
- For **regression**, typically a linear activation function (i.e.,  $\phi(z) = z$ ).
- For **classification**:

# Final Layer Activation Function

- The final layer activation function is chosen **depending on the task**.
- For **regression**, typically a linear activation function (i.e.,  $\phi(z) = z$ ).
- For **classification**:

## Binary Classification

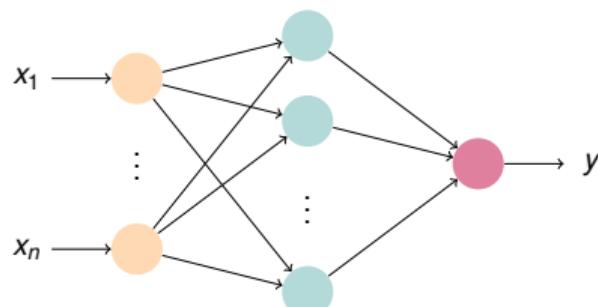


$$y = \text{sigmoid}(z)$$

# Final Layer Activation Function

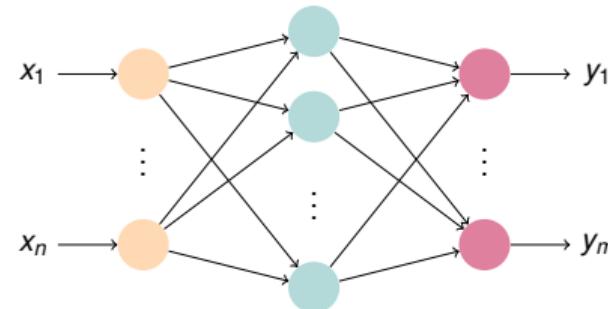
- The final layer activation function is chosen **depending on the task**.
- For **regression**, typically a linear activation function (i.e.,  $\phi(z) = z$ ).
- For **classification**:

Binary Classification



$$y = \text{sigmoid}(z)$$

Multi-class classification



$$y_i = \text{softmax}(\mathbf{z}) = \frac{e^{z_i}}{\sum_{k=1}^m e^{z_k}}$$

# Universal Approximation Theorem

**Theorem 1.** *Let  $\sigma$  be any continuous discriminatory function. Then finite sums of the form*

$$G(x) = \sum_{j=1}^N \alpha_j \sigma(y_j^T x + \theta_j) \quad (2)$$

*are dense in  $C(I_n)$ . In other words, given any  $f \in C(I_n)$  and  $\varepsilon > 0$ , there is a sum,  $G(x)$ , of the above form, for which*

$$|G(x) - f(x)| < \varepsilon \quad \text{for all } x \in I_n.$$

# Universal Approximation Theorem

**Theorem 1.** *Let  $\sigma$  be any continuous discriminatory function. Then finite sums of the form*

$$G(x) = \sum_{j=1}^N \alpha_j \sigma(y_j^T x + \theta_j) \quad (2)$$

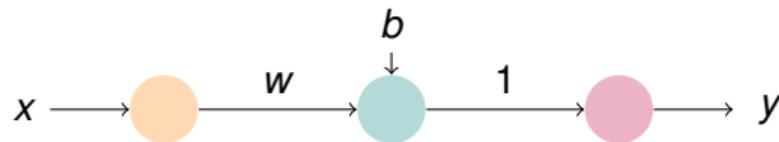
*are dense in  $C(I_n)$ . In other words, given any  $f \in C(I_n)$  and  $\varepsilon > 0$ , there is a sum,  $G(x)$ , of the above form, for which*

$$|G(x) - f(x)| < \varepsilon \quad \text{for all } x \in I_n.$$

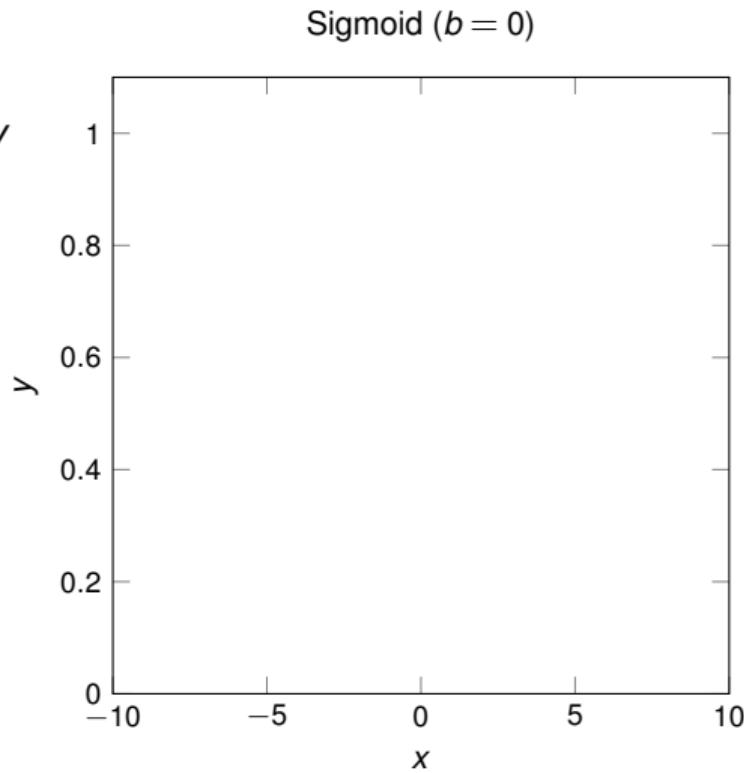
Huh?!?

G. Cybenko "Approximation by superpositions of a sigmoidal function," Mathematics of Control, Signals and Systems, 1989.

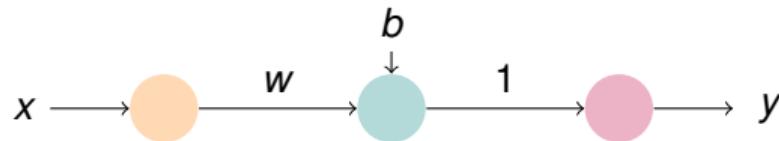
# Neural Network as a Step Function



$$y = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

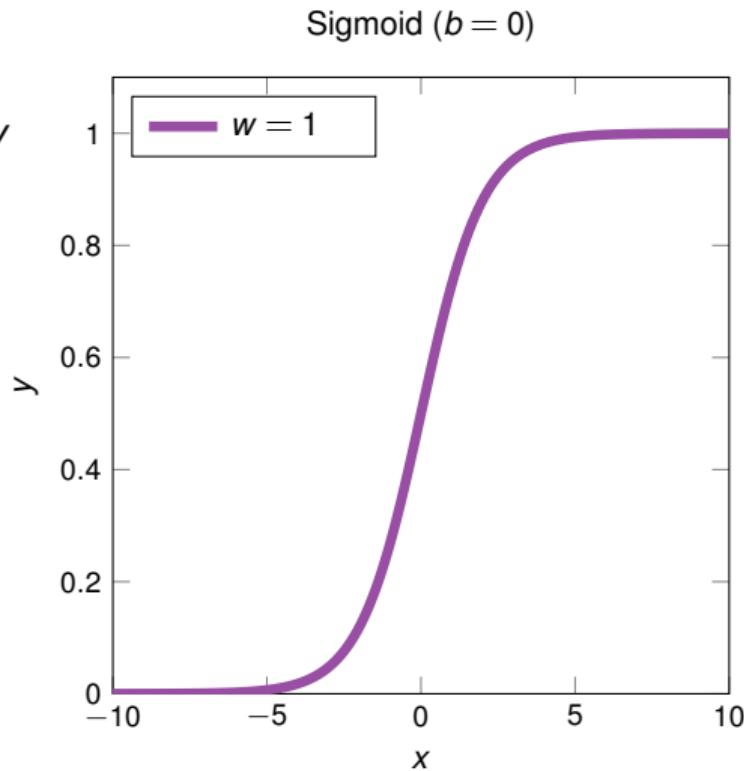


# Neural Network as a Step Function

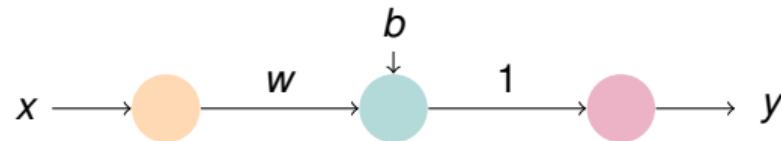


$$y = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

1. Assume  $b = 0$  and vary  $w$

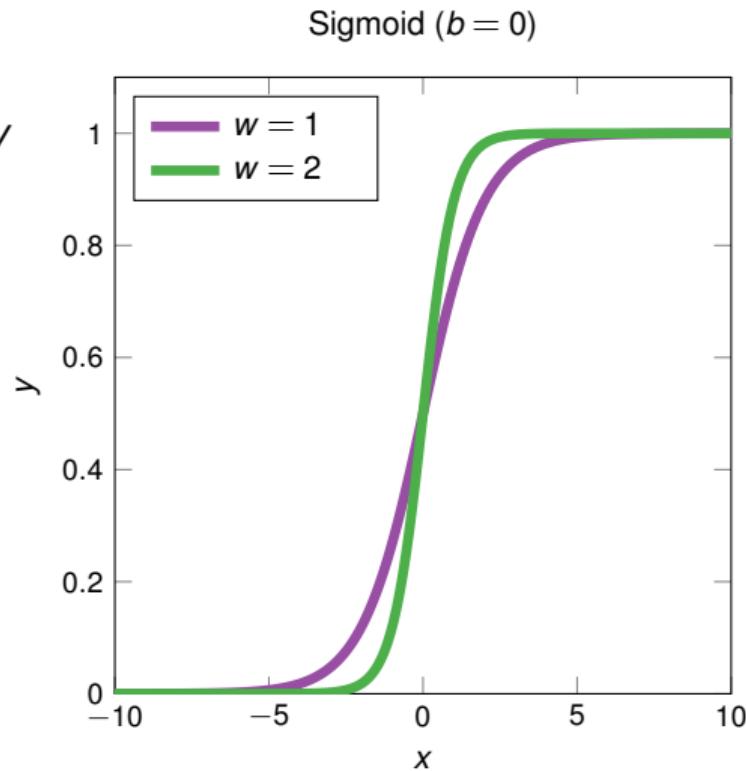


# Neural Network as a Step Function

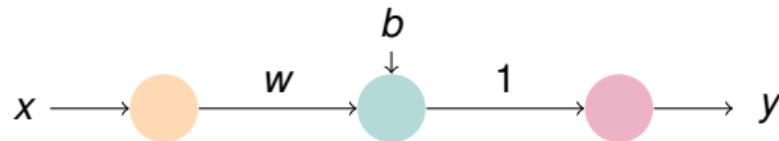


$$y = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

1. Assume  $b = 0$  and vary  $w$

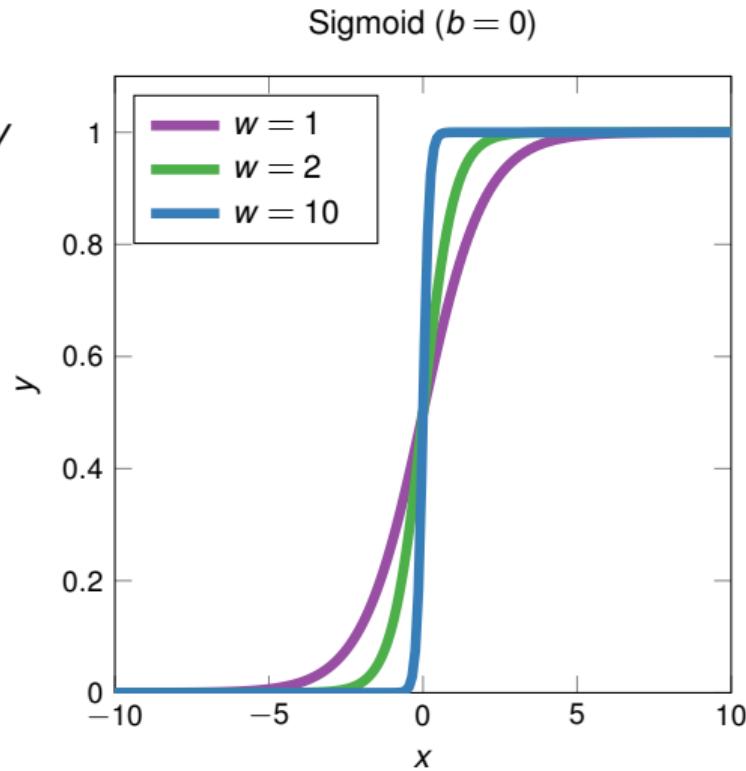


# Neural Network as a Step Function

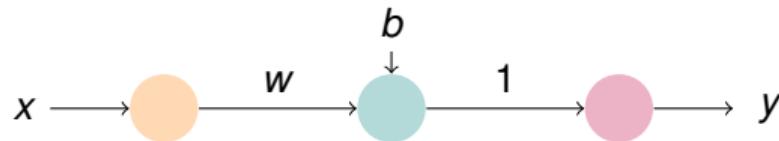


$$y = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

1. Assume  $b = 0$  and vary  $w$

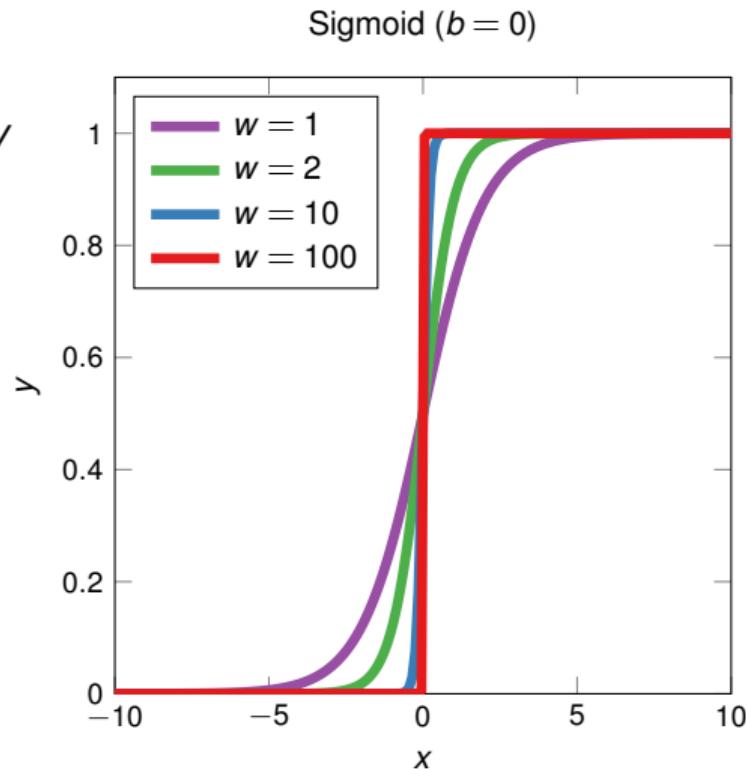


# Neural Network as a Step Function

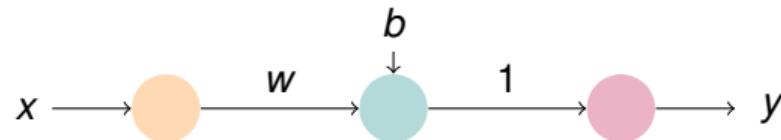


$$y = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

1. Assume  $b = 0$  and vary  $w$

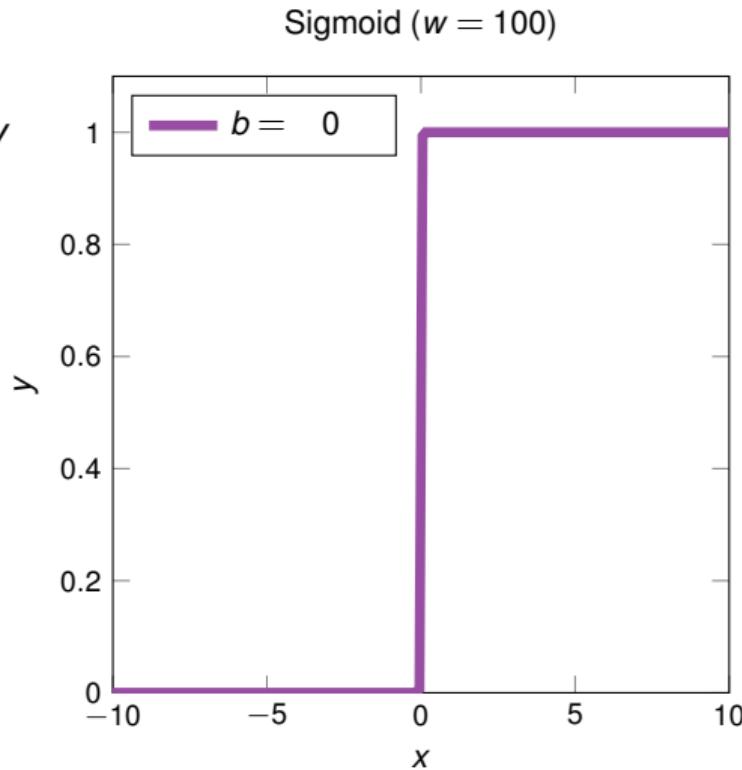


# Neural Network as a Step Function

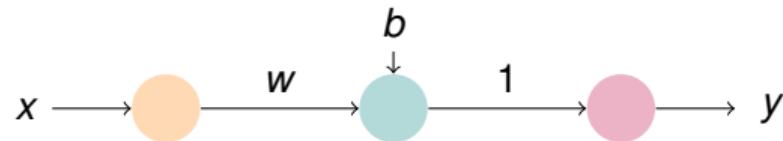


$$y = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

1. Assume  $b = 0$  and vary  $w$
2. Assume  $w = 100$  and vary  $b$

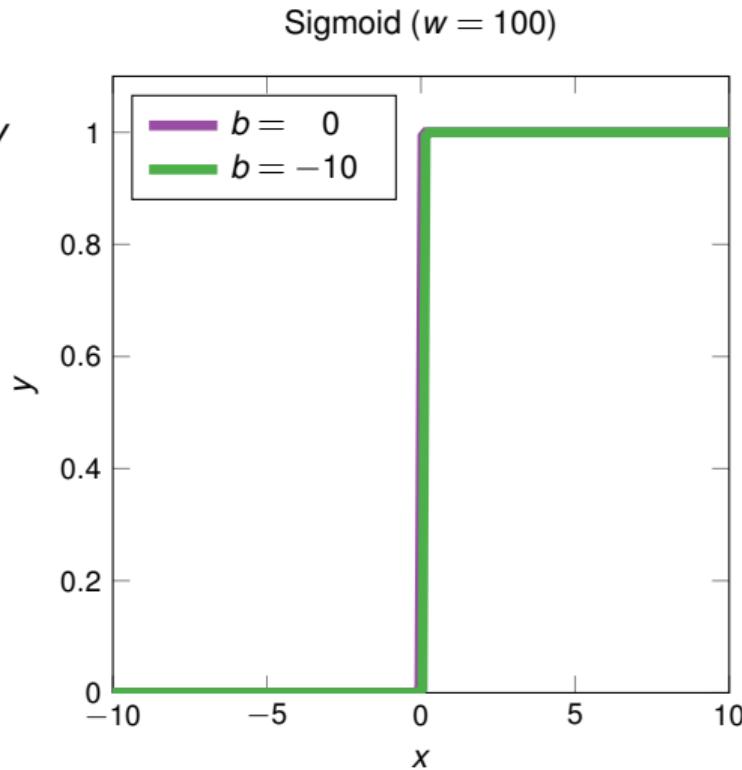


# Neural Network as a Step Function

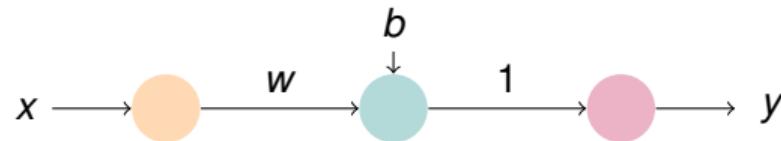


$$y = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

1. Assume  $b = 0$  and vary  $w$
2. Assume  $w = 100$  and vary  $b$

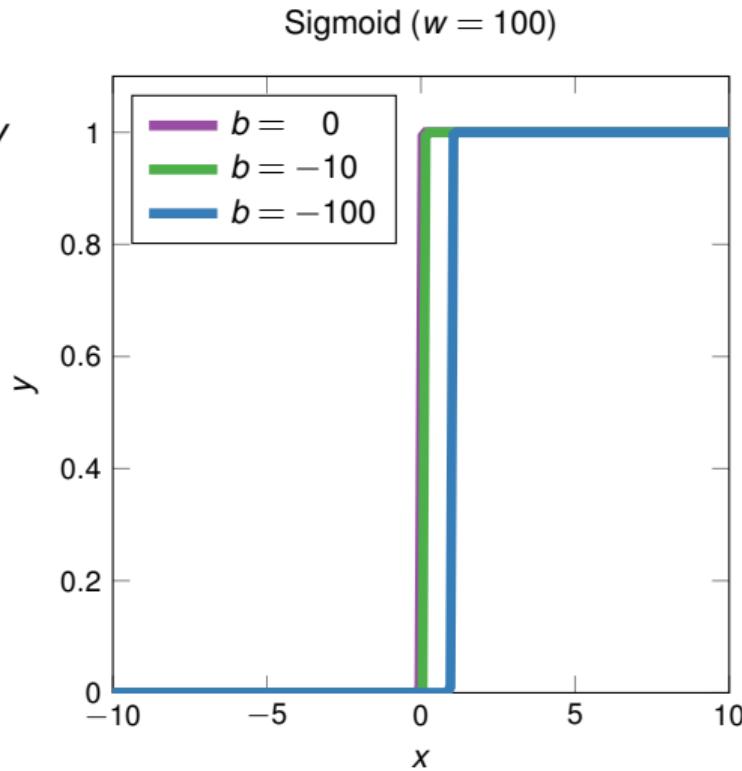


# Neural Network as a Step Function

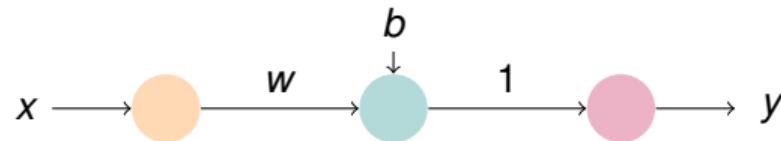


$$y = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

1. Assume  $b = 0$  and vary  $w$
2. Assume  $w = 100$  and vary  $b$

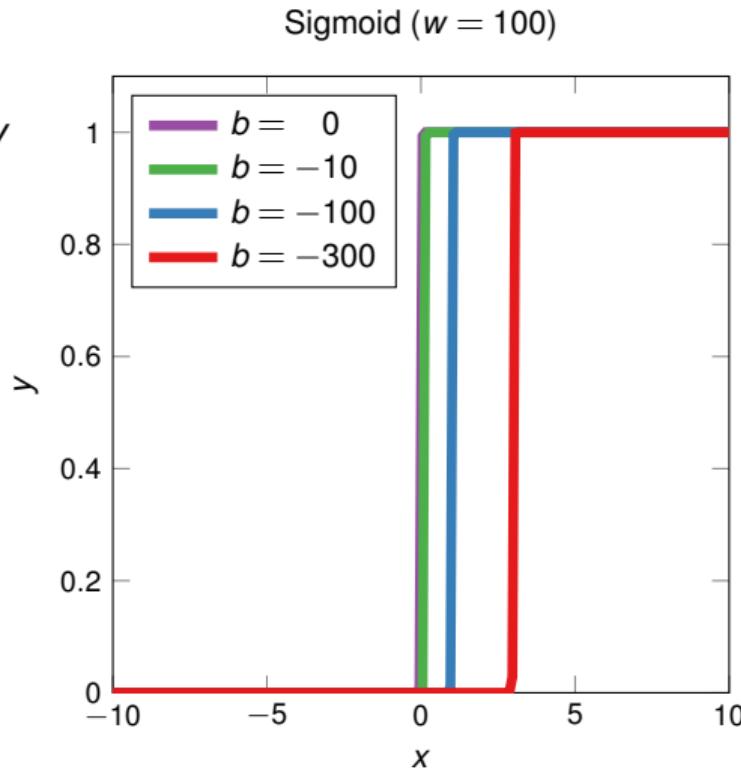


# Neural Network as a Step Function

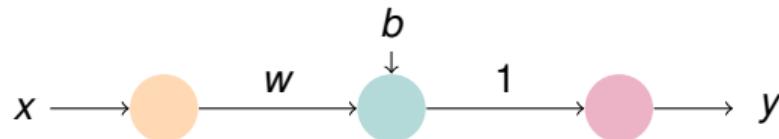


$$y = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

1. Assume  $b = 0$  and vary  $w$
2. Assume  $w = 100$  and vary  $b$



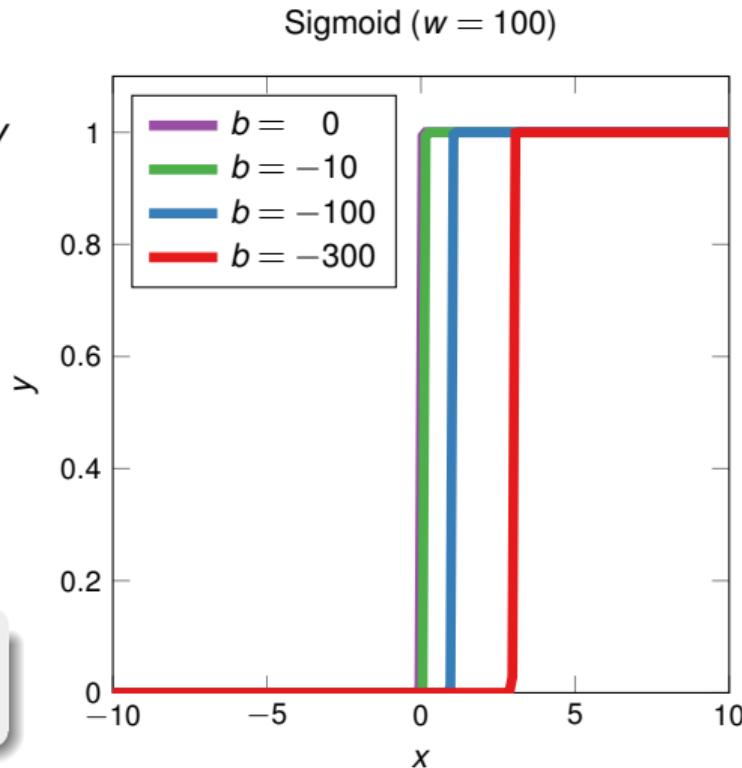
# Neural Network as a Step Function



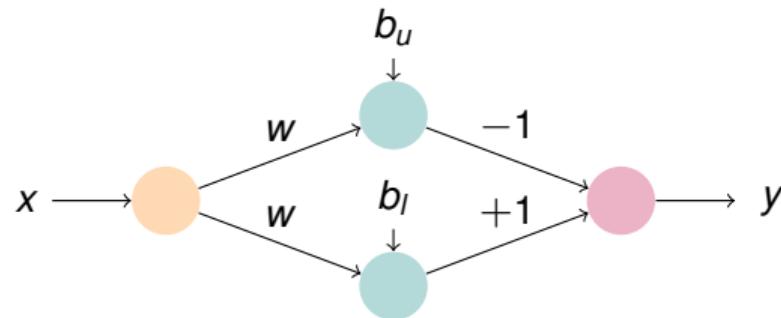
$$y = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

1. Assume  $b = 0$  and vary  $w$
2. Assume  $w = 100$  and vary  $b$

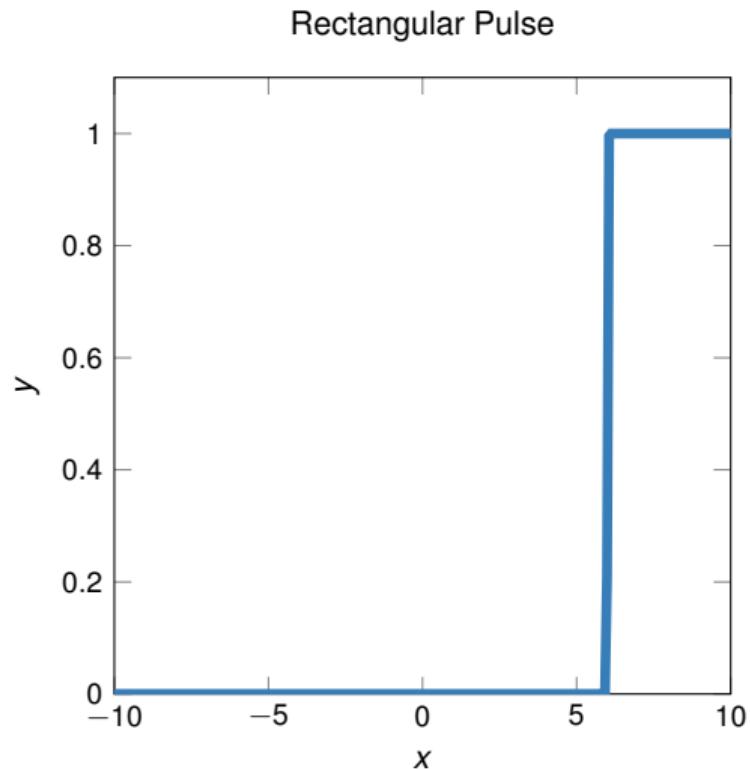
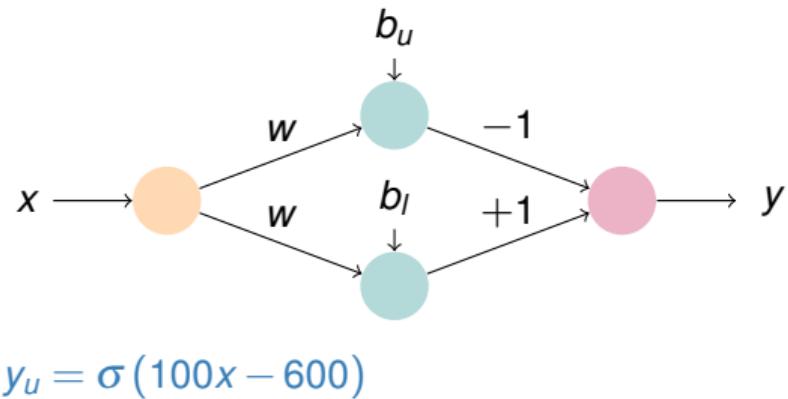
In general,  $y = \sigma(wx + b)$  approximates a **step function** at  $x = -\frac{b}{w}$ .



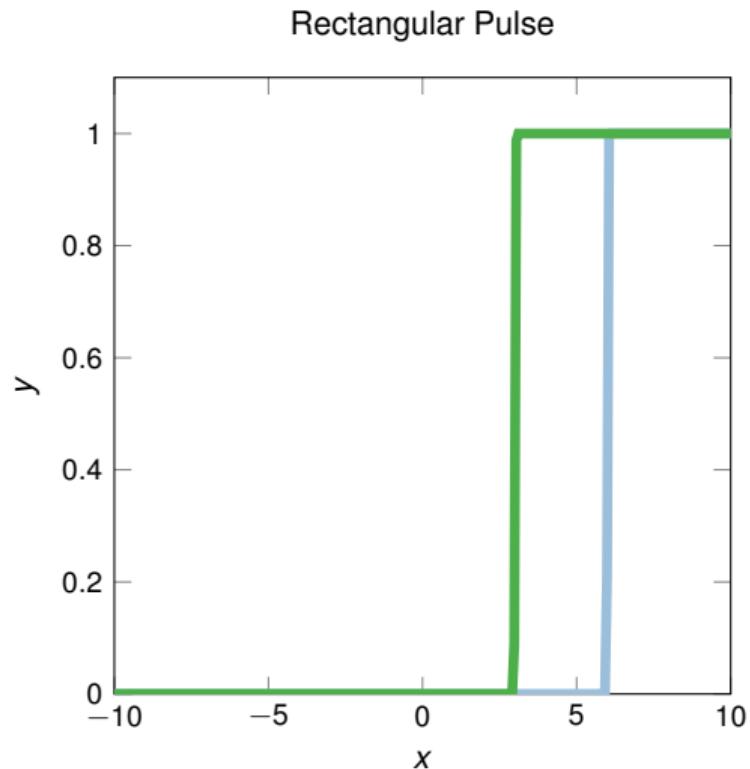
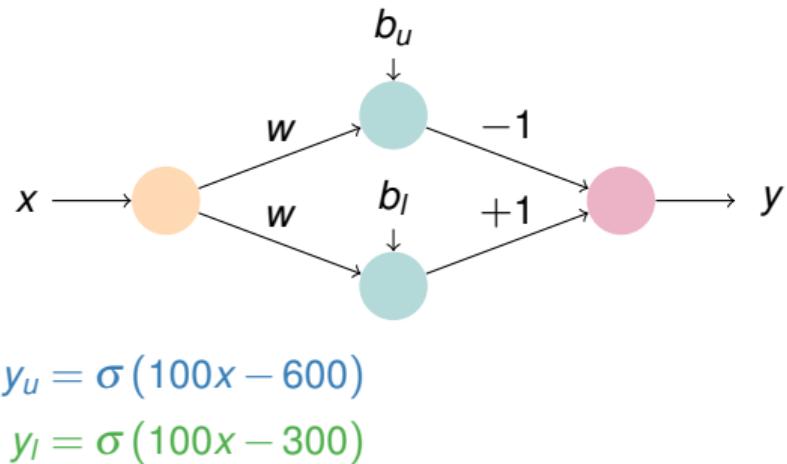
# Neural Network as a Rectangular Function



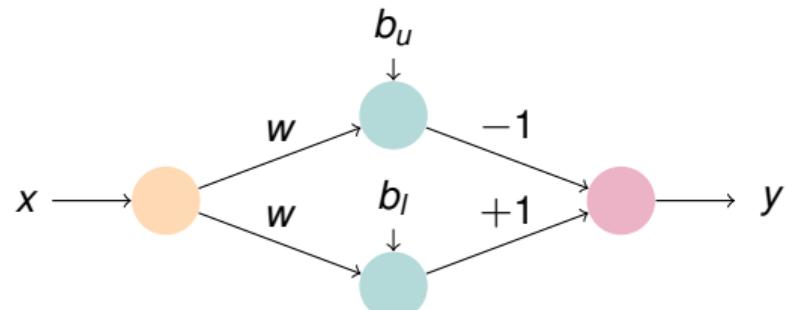
# Neural Network as a Rectangular Function



# Neural Network as a Rectangular Function



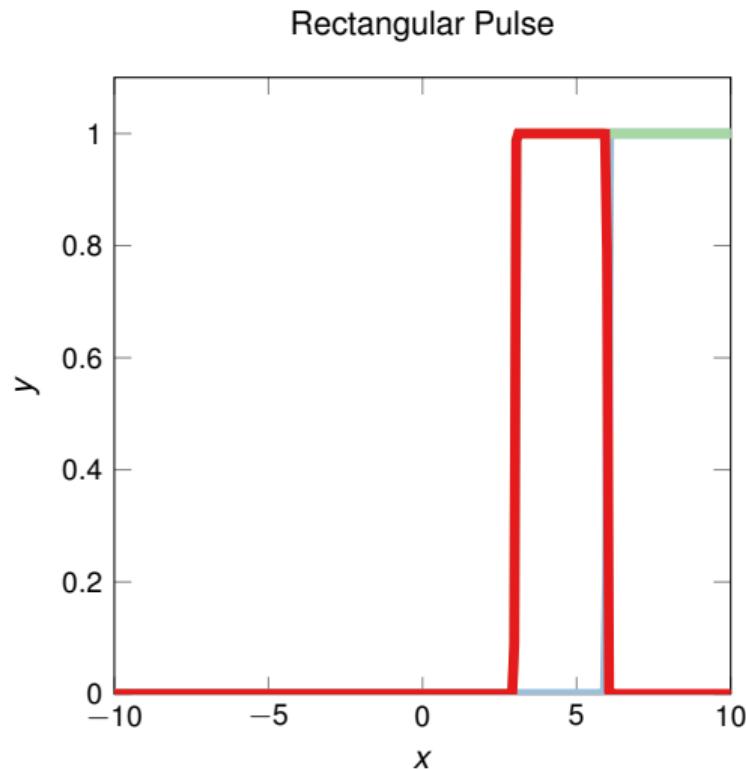
# Neural Network as a Rectangular Function



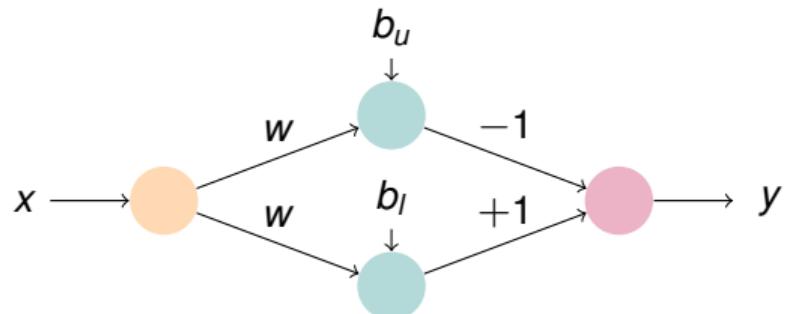
$$y_u = \sigma(100x - 600)$$

$$y_l = \sigma(100x - 300)$$

$$y = y_l - y_u = \sigma(100x - 300) - \sigma(100x - 600)$$



# Neural Network as a Rectangular Function

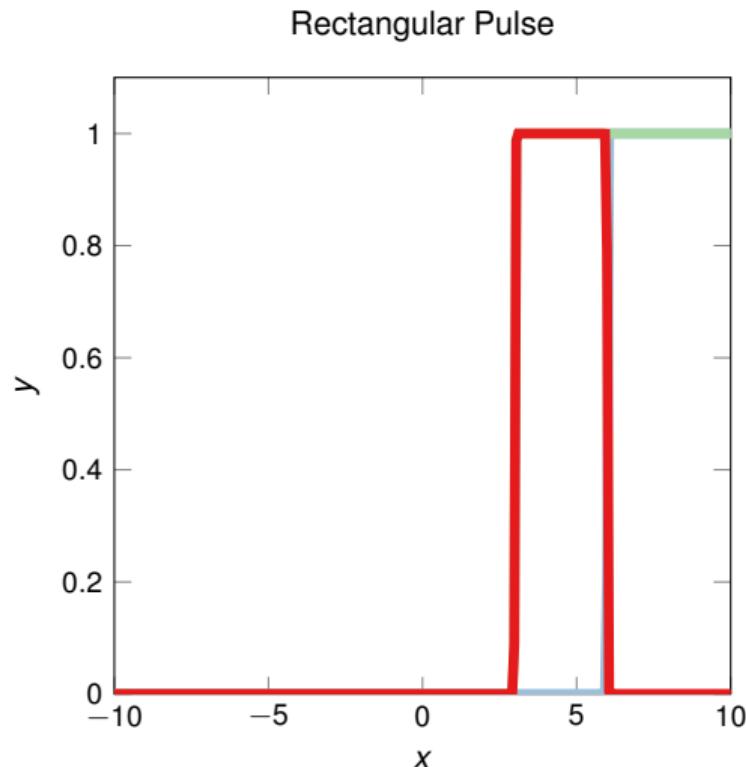


$$y_u = \sigma(100x - 600)$$

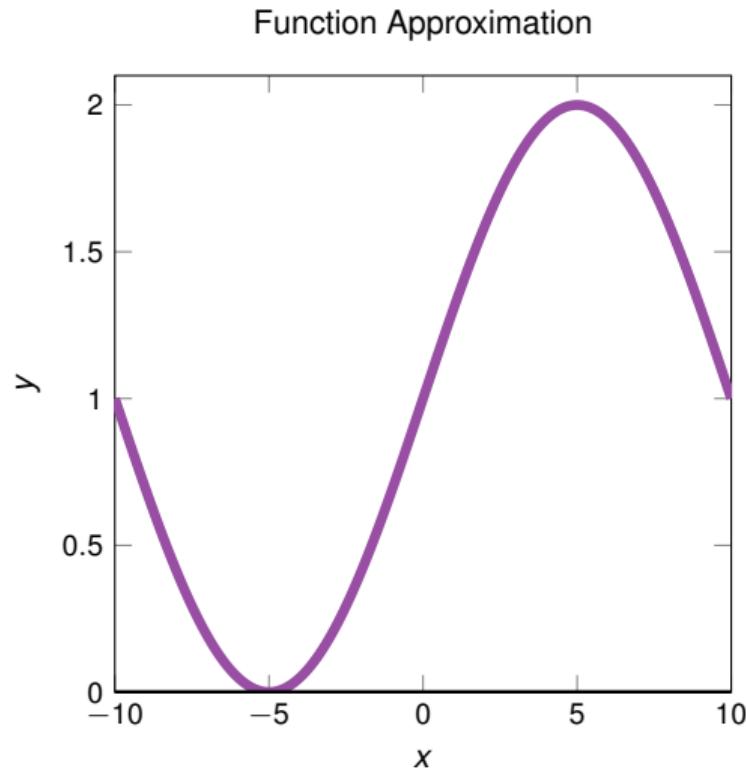
$$y_l = \sigma(100x - 300)$$

$$y = y_l - y_u = \sigma(100x - 300) - \sigma(100x - 600)$$

In general,  $y = \sigma(wx + b_l) - \sigma(wx + b_u)$  approximates a **rectangular function** from  $x_l = -\frac{b_l}{w}$  to  $x_u = -\frac{b_u}{w}$ .



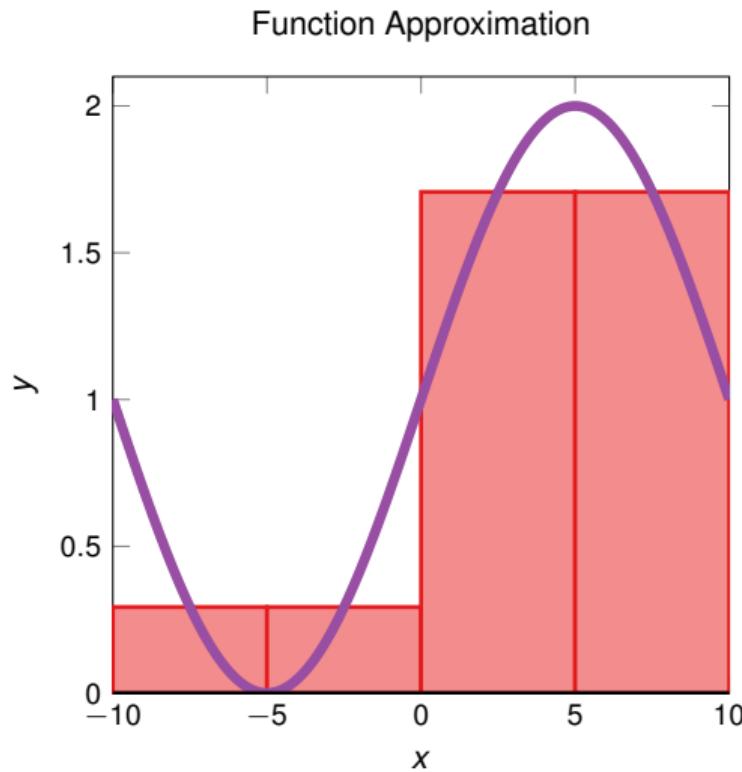
# Neural Network as a Universal Approximator



# Neural Network as a Universal Approximator

**Generalization:**

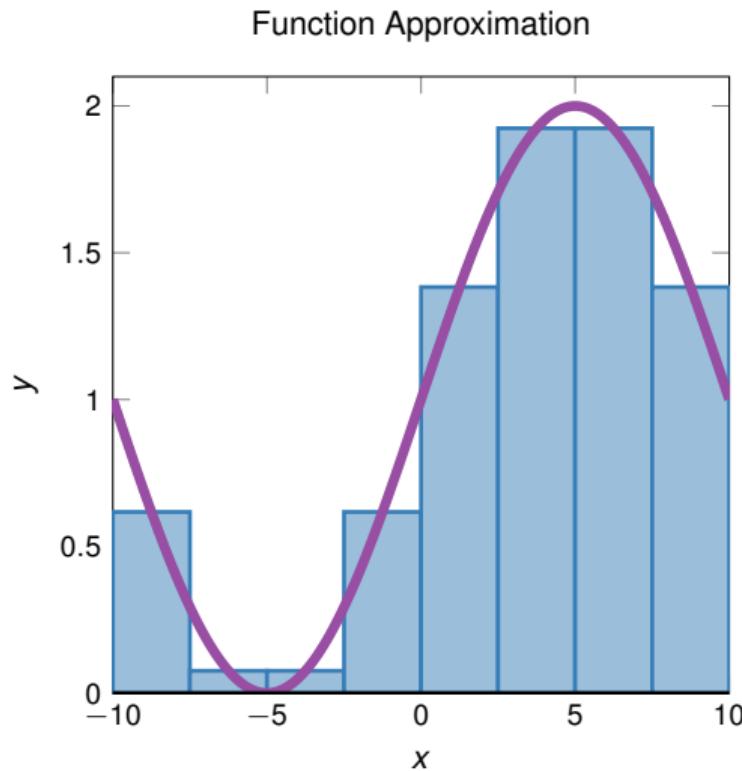
$$y = \sum_{j=1}^4 \alpha_j (\sigma(wx - b_{l,j}) - \sigma(wx - b_{u,j}))$$



# Neural Network as a Universal Approximator

**Generalization:**

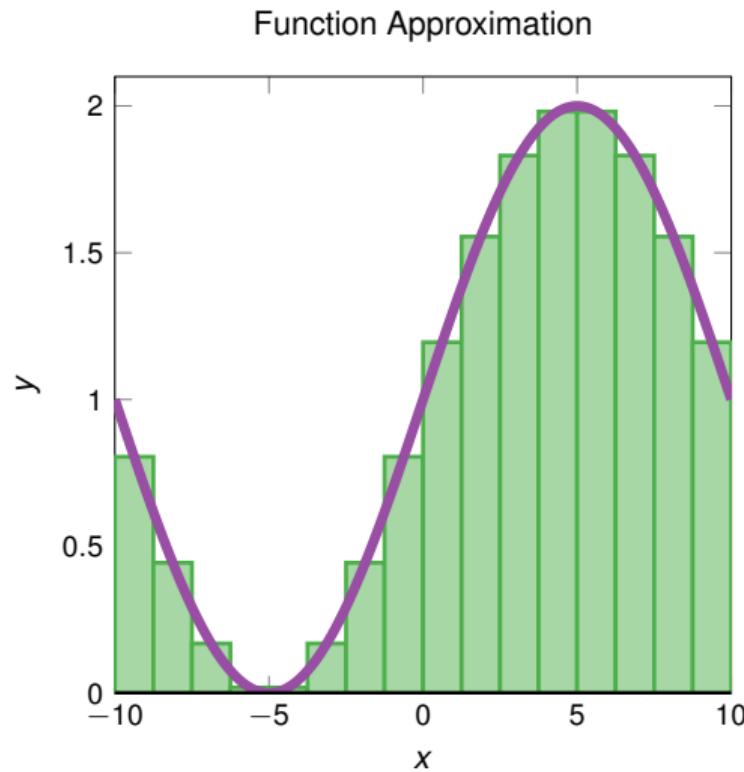
$$y = \sum_{j=1}^8 \alpha_j (\sigma(wx - b_{l,j}) - \sigma(wx - b_{u,j}))$$



# Neural Network as a Universal Approximator

**Generalization:**

$$y = \sum_{j=1}^{16} \alpha_j (\sigma(wx - b_{l,j}) - \sigma(wx - b_{u,j}))$$



# Neural Network as a Universal Approximator

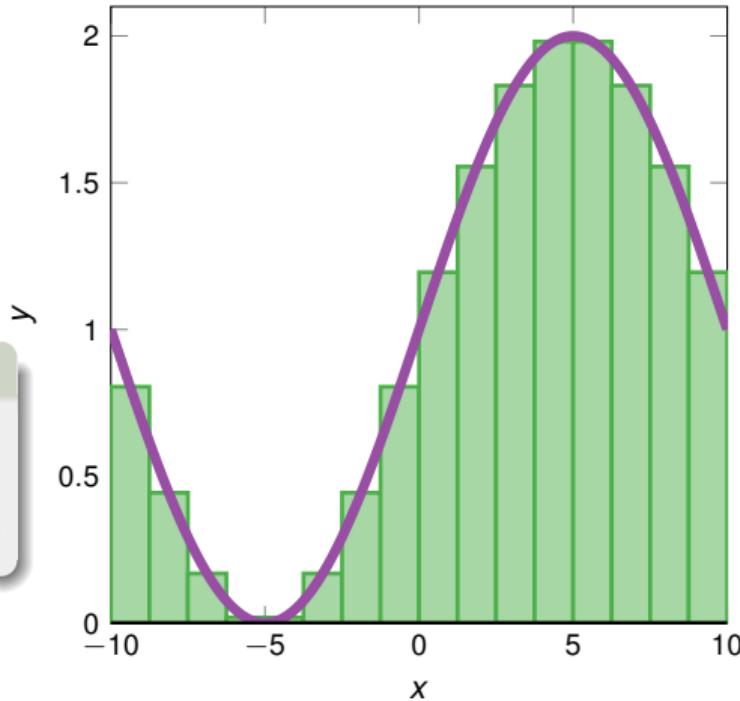
## Generalization:

$$y = \sum_{j=1}^{16} \alpha_j (\sigma(wx - b_{l,j}) - \sigma(wx - b_{u,j}))$$

### Result

We can use  $N$  rectangular functions constructed from  $2N$  neurons to **approximate any function arbitrarily well!**

Function Approximation



# Conclusion

## Summary:

1. Neural networks (NNs) have a long history, but recently became extremely popular.
2. NNs consist of artificial neurons that are brain-inspired.
3. NNs perform matrix-vector multiplications, vector additions, and activation functions.
4. NNs are universal function approximators.

# Conclusion

## Summary:

1. Neural networks (NNs) have a long history, but recently became extremely popular.
2. NNs consist of artificial neurons that are brain-inspired.
3. NNs perform matrix-vector multiplications, vector additions, and activation functions.
4. NNs are universal function approximators.

## Next time:

- **NN training:** cost functions, (improved) gradient descent, backpropagation, weight initialization, data normalization.